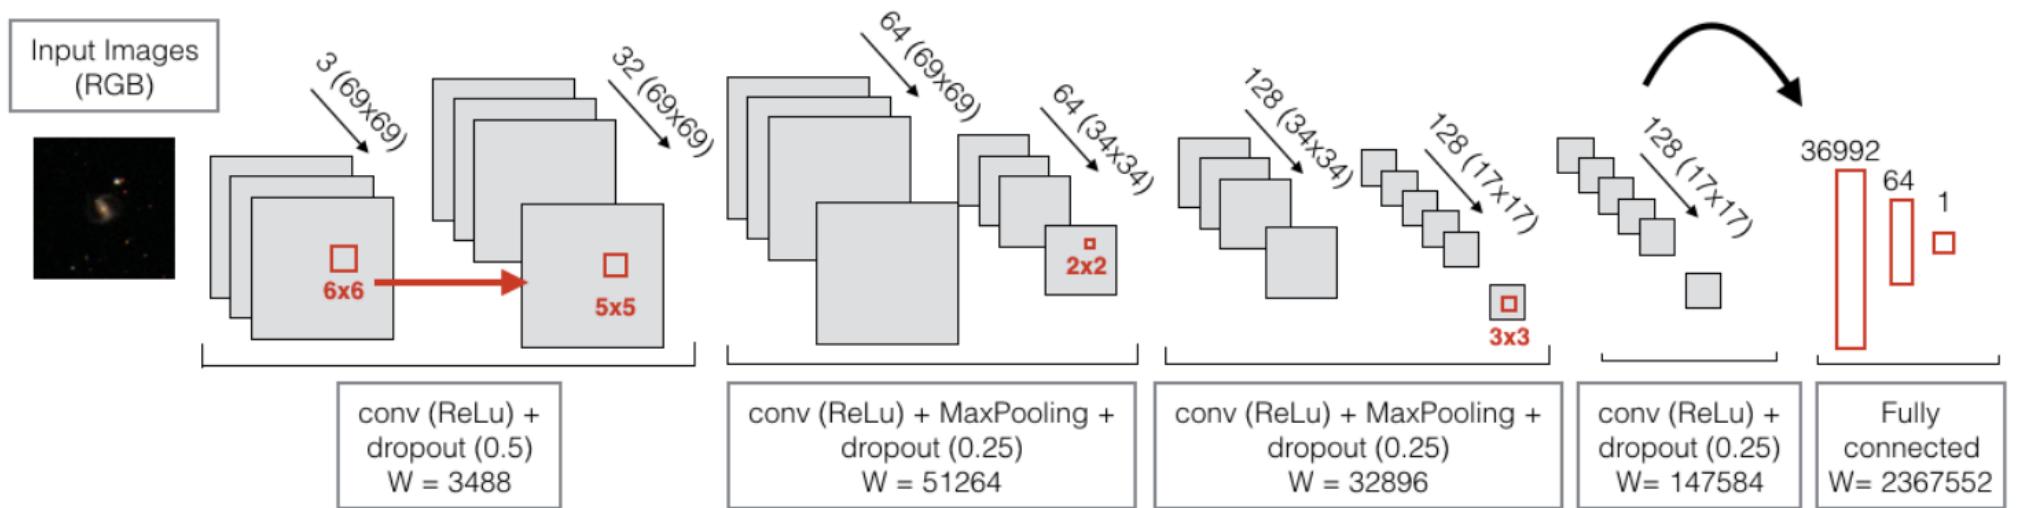


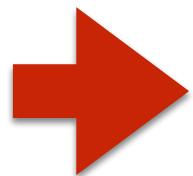
## PART IV: BEYOND CLASSIFICATION: IMAGE2IMAGE NETWORKS

# UP TO NOW CNNs MAP IMAGES (SIGNALS) INTO FLOATS



Dominguez-Sanchez+18

Classification has its limits ....



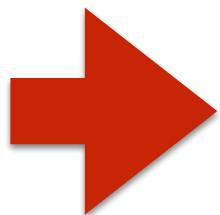
**HOW DO I CLASSIFY THIS IMAGE?**

Classification has its limits ....



classification

person, sheep, dog



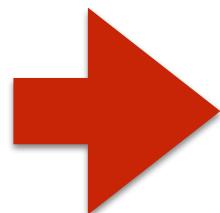
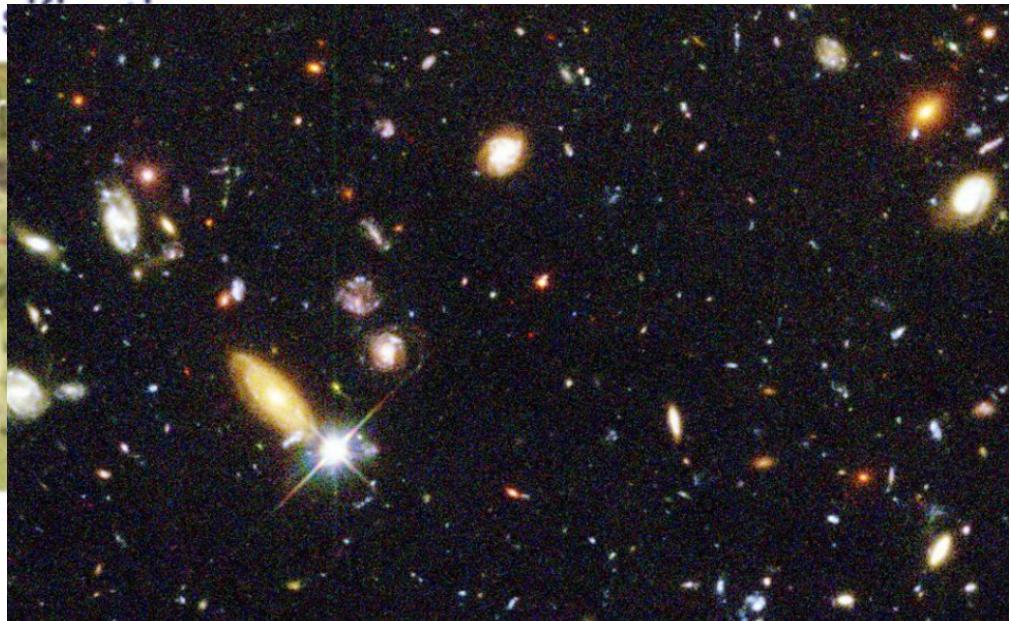
**HOW DO I CLASSIFY THIS IMAGE?**

Classification has its limits ....



classifi

per



**HOW DO I CLASSIFY THIS IMAGE?**

# Going beyond classification: increasing complexity

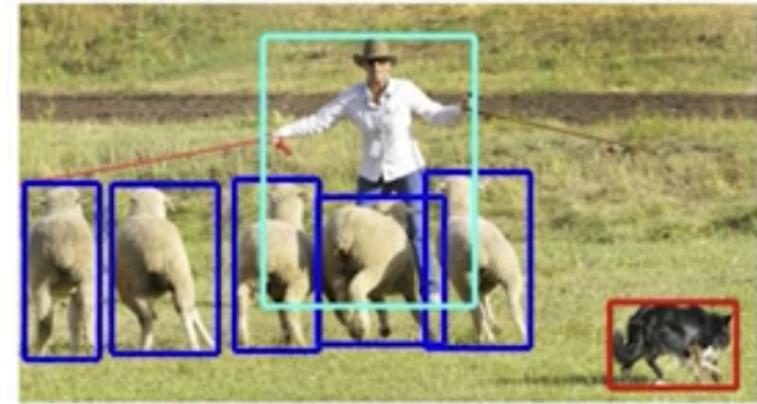
classification



semantic segmentation



object detection

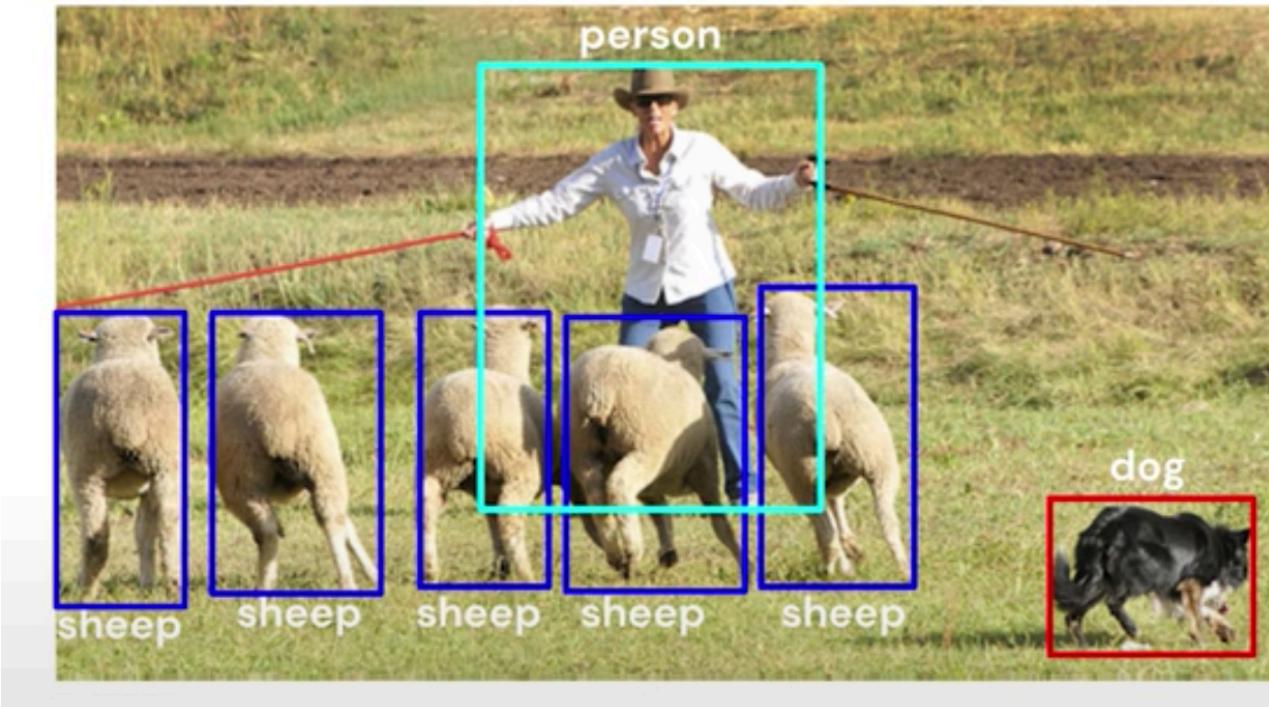


instance segmentation



# Object detection

First task is to find a bounding box for every object. How we do that?



## Inputs

- RGB image  $H \times W \times 3$

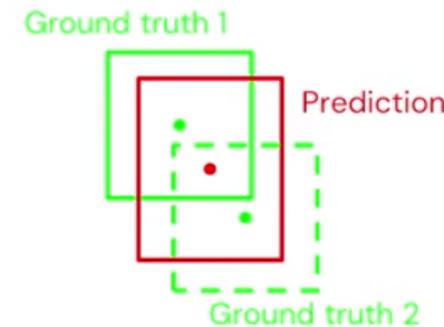
## Targets

- Class label one\_hot  $0\ 0\ 0\ 1\ 0\ ...$
- Object bounding box  
 $(x_c, y_c, h, w)$

for all the objects present in the scene

# WHAT WOULD BE THE LOSS FUNCTION OF SUCH A PROBLEM?

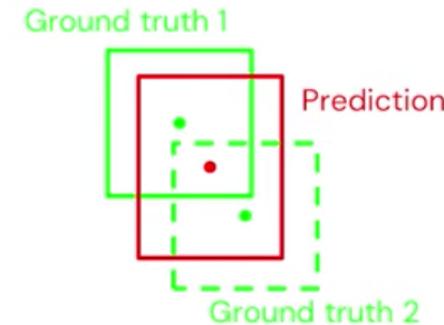
$$\frac{1}{N} \sum_{i=1}^N (y_i - p_i)^2$$



IT IS A REGRESSION WITH A SIMPLE QUADRATIC LOSS.  
WE TRY TO FIND THE BEST COORDINATES OF THE  
BOUNDING BOX.

# WHAT WOULD BE THE LOSS FUNCTION OF SUCH A PROBLEM?

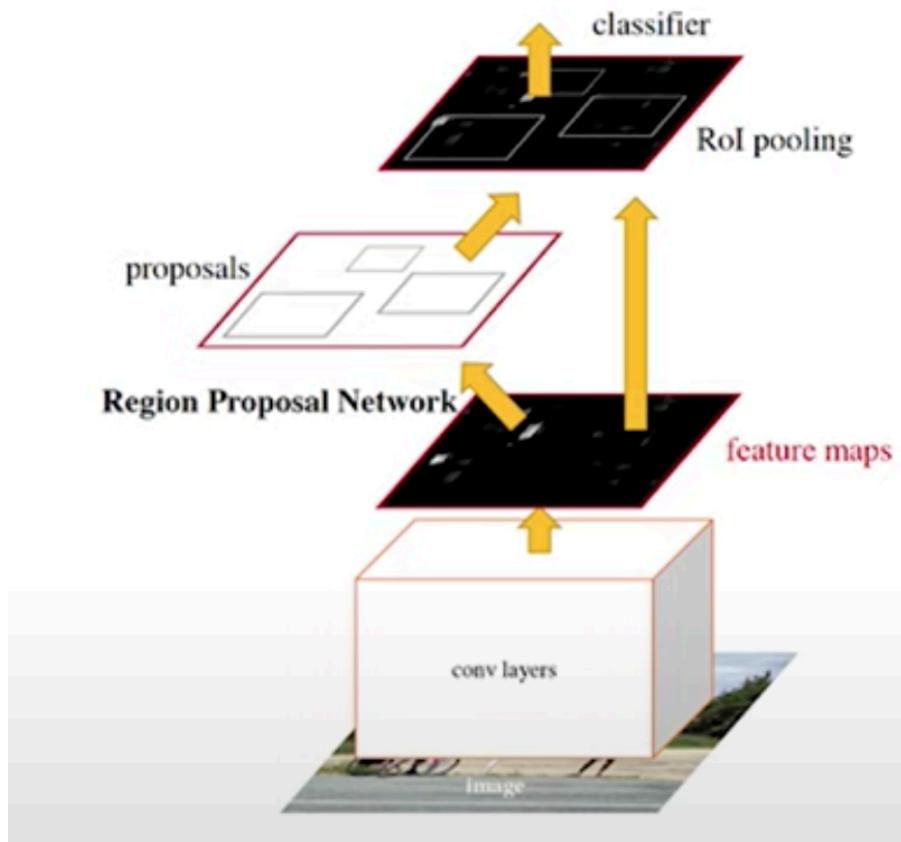
$$\frac{1}{N} \sum_{i=1}^N (y_i - p_i)^2$$



IT IS A REGRESSION WITH A SIMPLE QUADRATIC LOSS.  
WE TRY TO FIND THE BEST COORDINATES OF THE  
BOUNDING BOX.

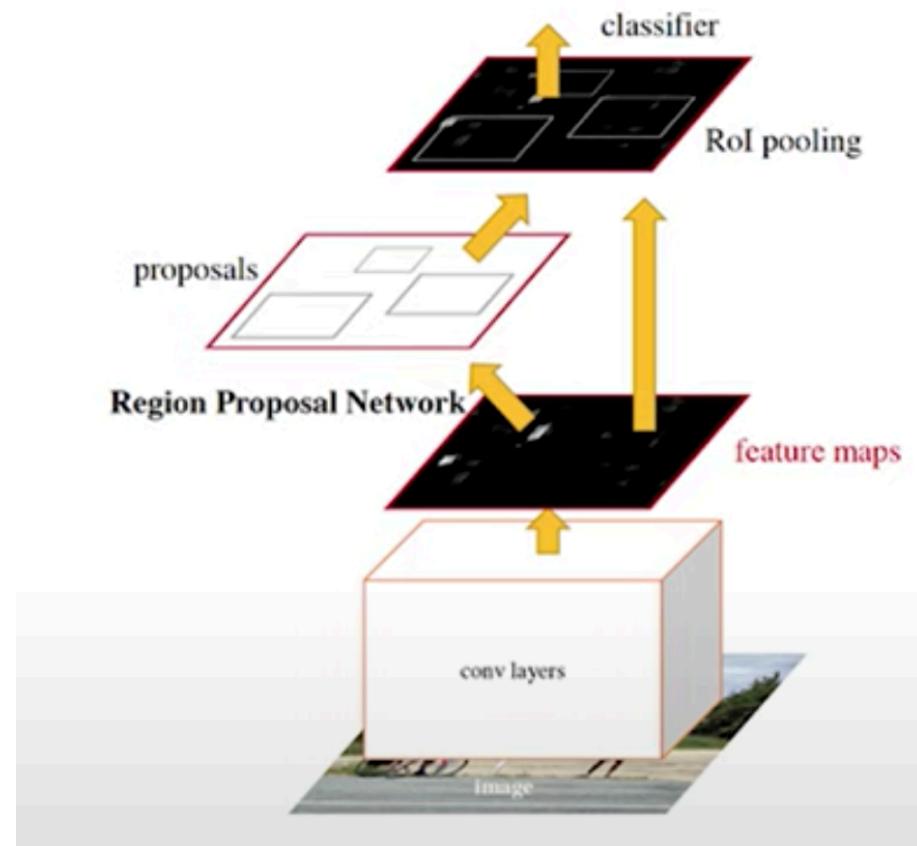
IT BECOMES MESSY QUITE RAPIDLY...

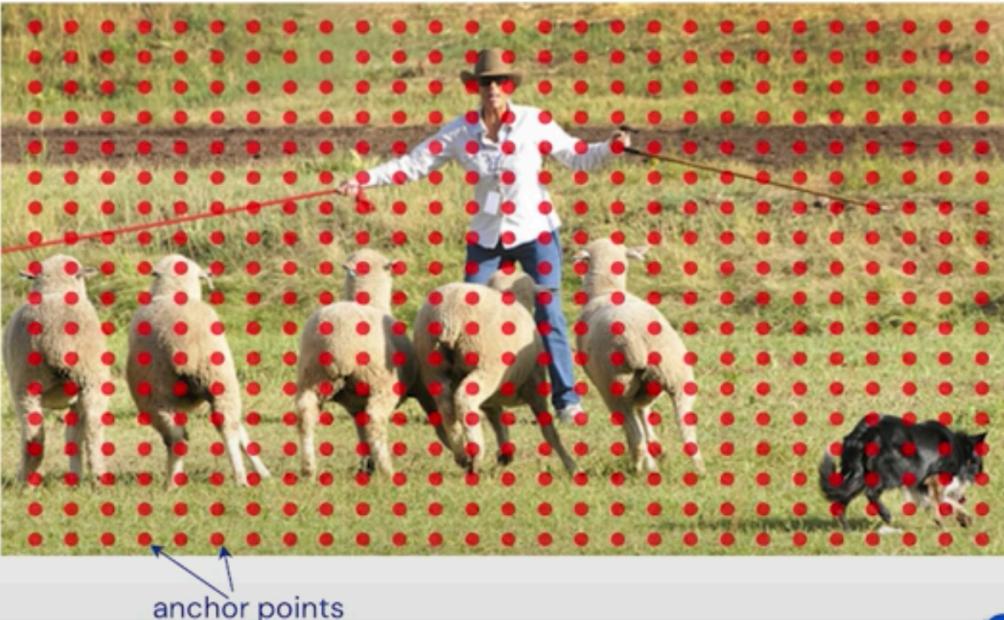
# EXAMPLE: FASTER R-CNN



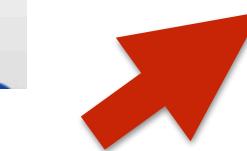
We divide the task in 2 steps:

1. Identify bounding box candidates (CLASSIFICATION)
2. Classify and Refine (REGRESSION)





Discretize Bounding Box Space



Choose n candidates per position

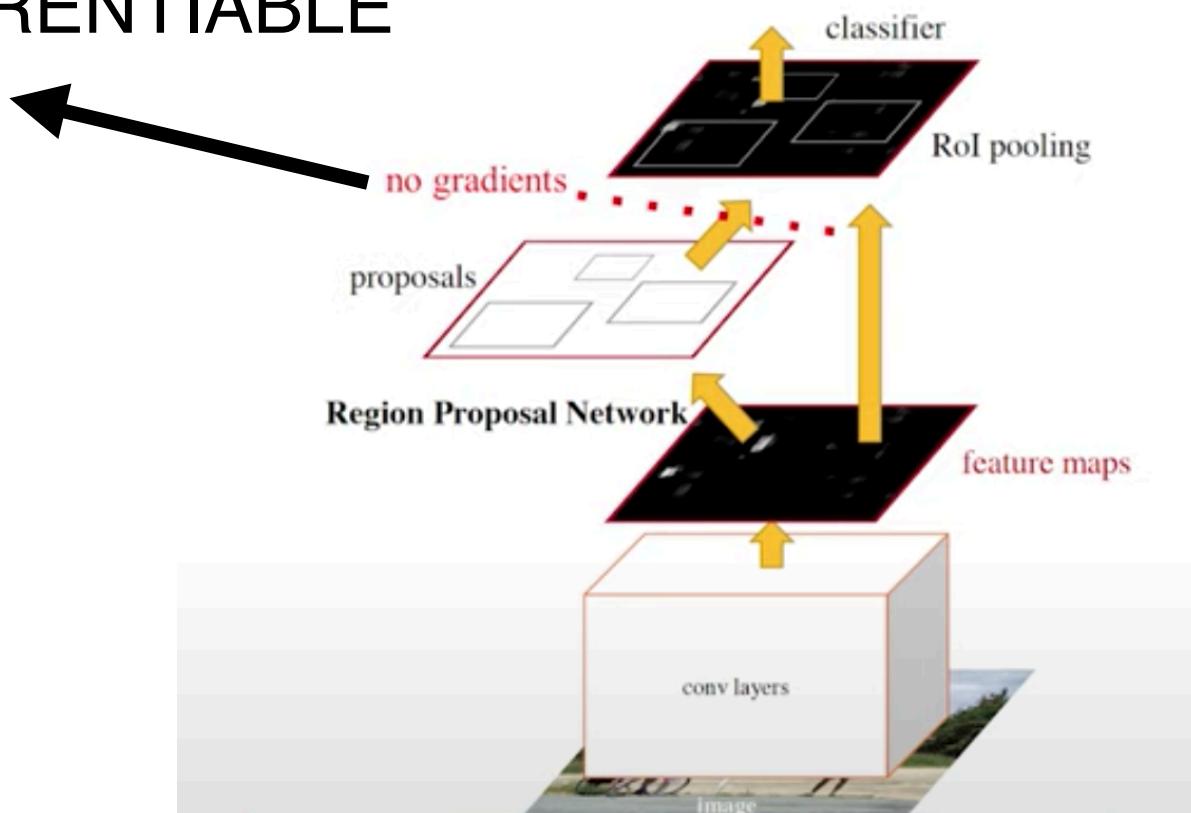


Predict objectness score (classification)

Sort and keep top K

# HIS IS NOT DIFFERENTIABLE

(FIXED)



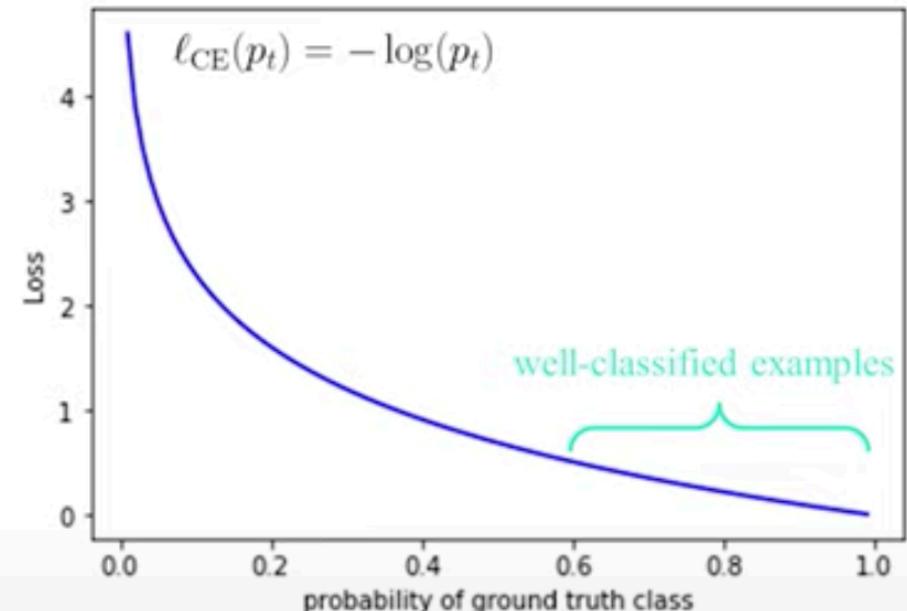
# WYS OF MAKING THIS DIFFERENTIABLE:

Spatial Transformer Networks (Jaderberg+18)

# WHY NOT DOING IT ONE STAGE?

MOST OF THE CANDIDATES  
ARE BACKGROUND,  
EASY TO IDENTIFY

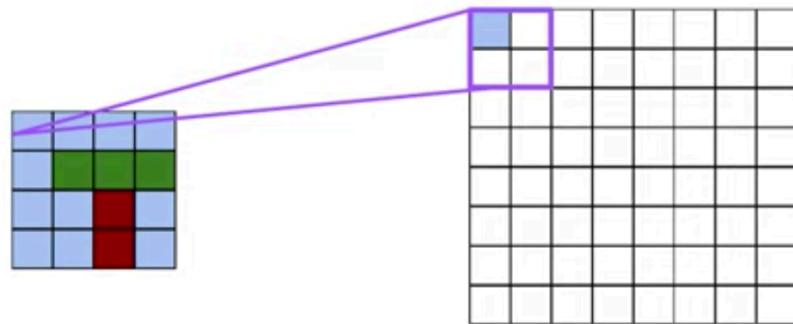
THE LOSS OF THE MANY  
EASY EXAMPLES  
DOMINATES OVER THE  
RARE USEFUL ONES



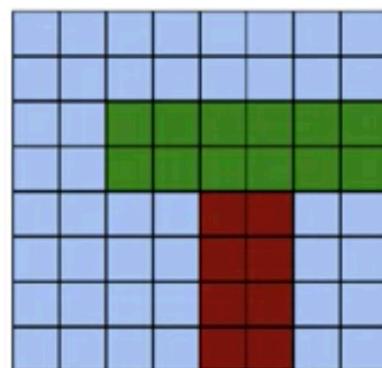
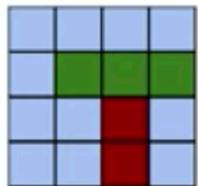
semantic segmentation



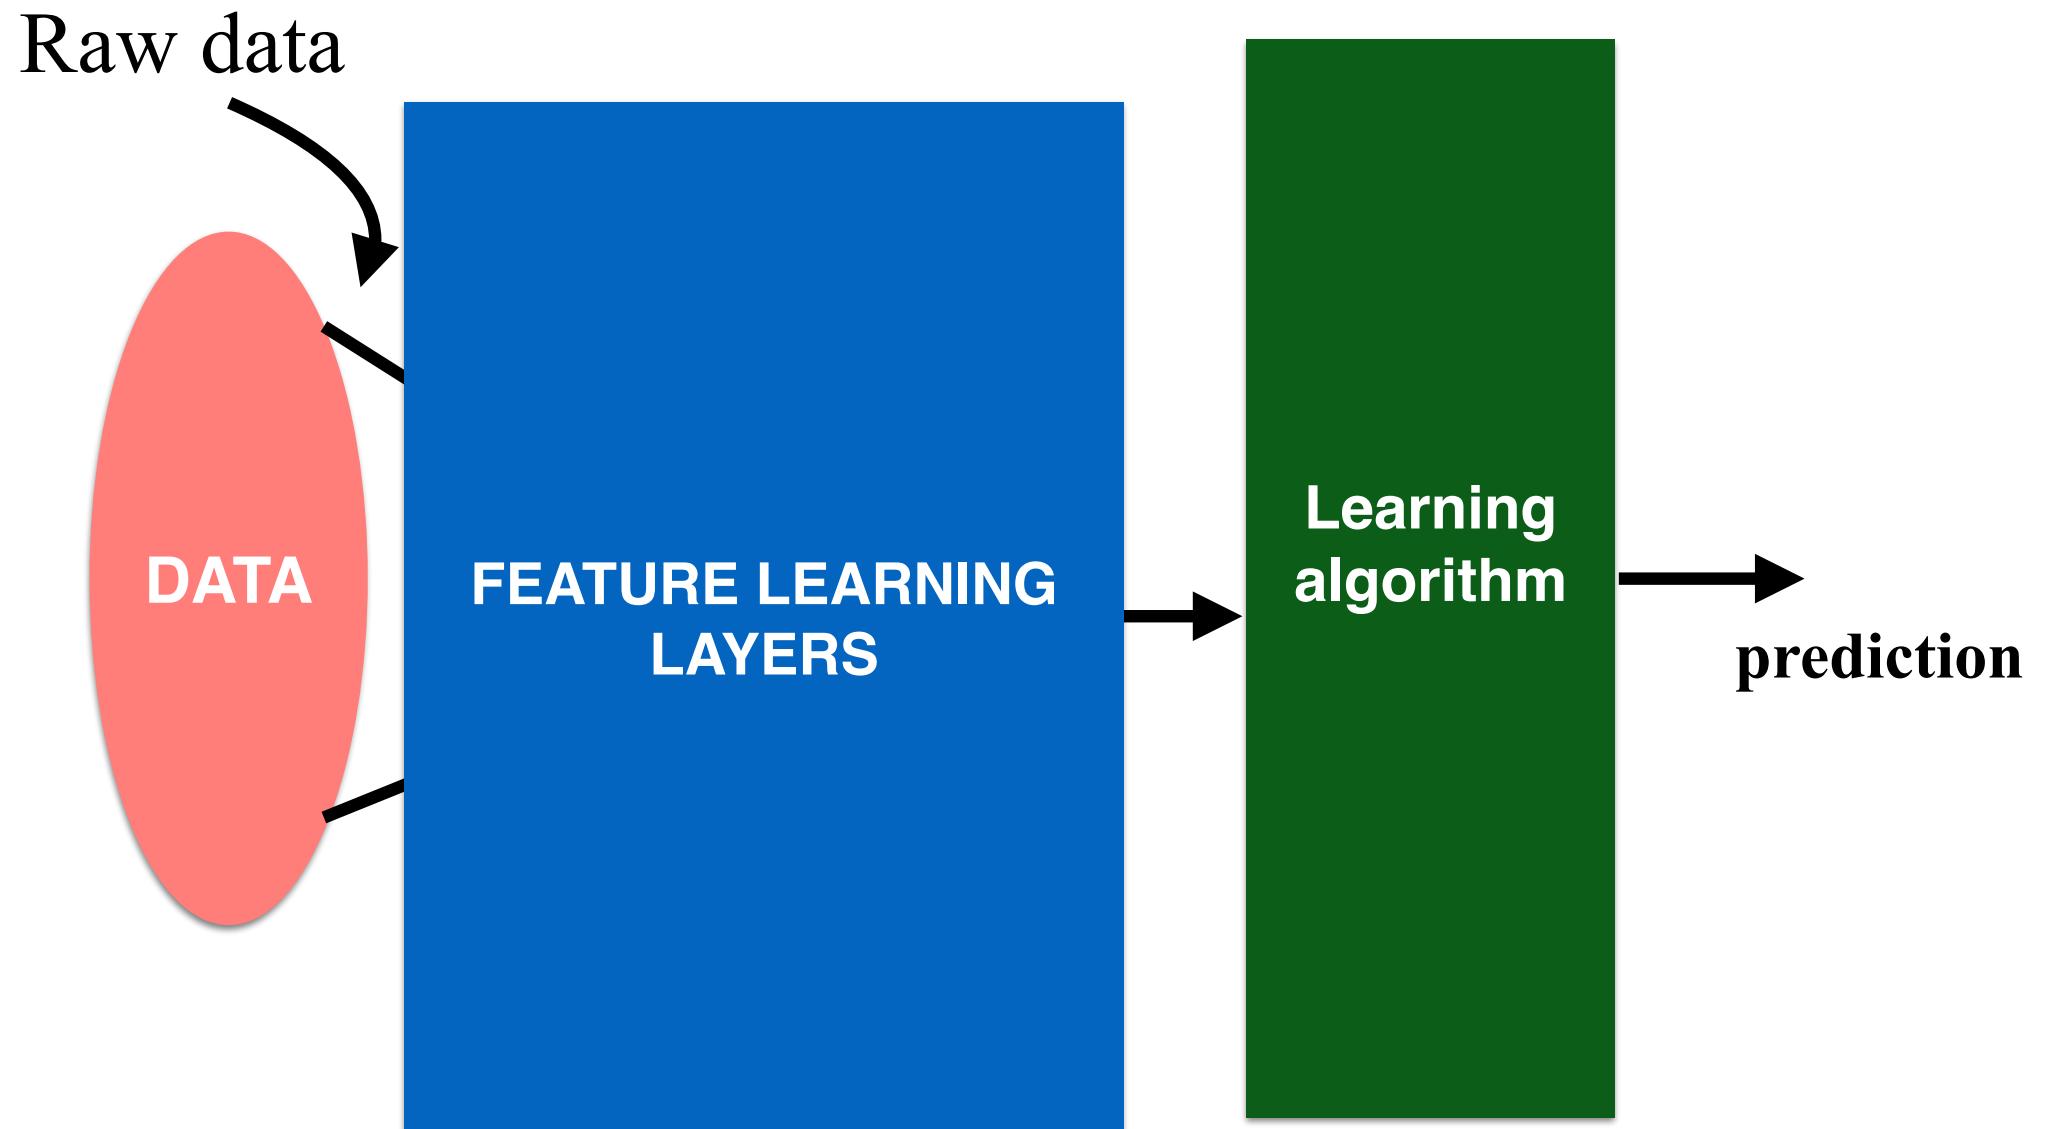
# UNPOOLING OPERATION (INVERSE OF POOLING)



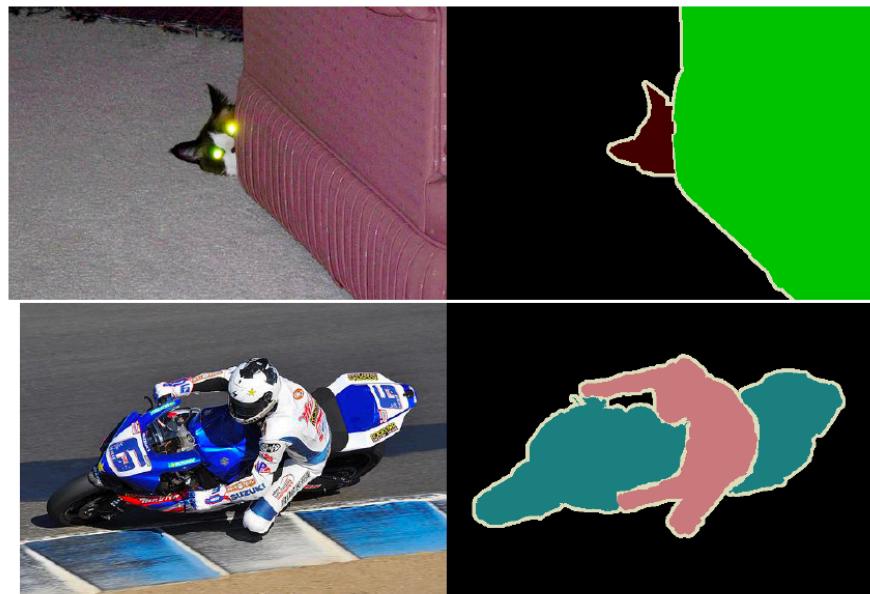
COPY PIXELS IN A  
GIVEN WINDOW



GENERATES  
LARGER IMAGES  
FROM SMALLER  
ONES

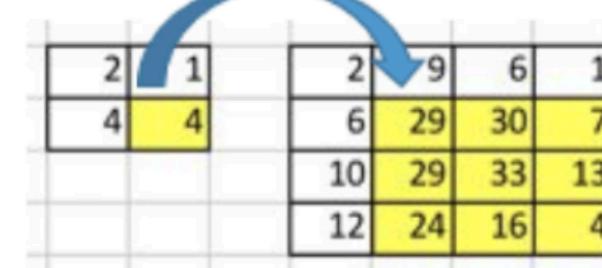
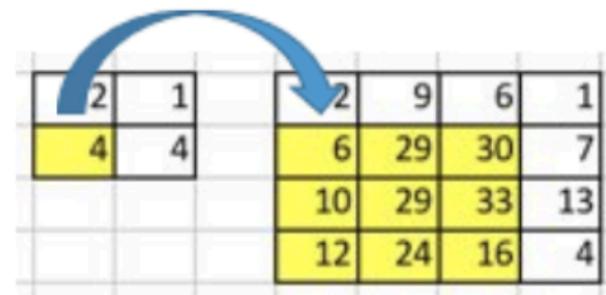
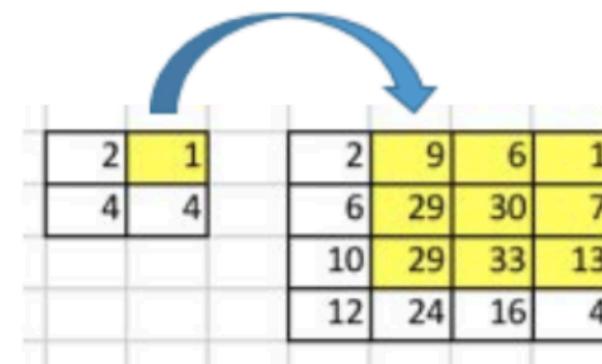
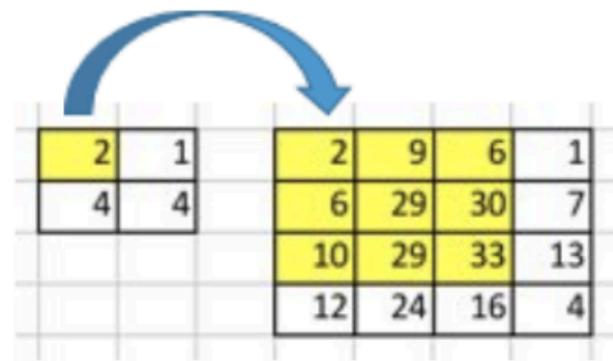


# FIRST: IMAGE SEGMENTATION WITH ENCODERS-DECODERS



# TRANSPOSED CONVOLUTION

ALLOWS TO INCREASE THE SIZE



Going Backward of Convolution

EXAMPLE TAKEN FROM HERE

# CONVOLUTION MATRIX

	0	1	2
0	1	4	1
1	1	4	3
2	3	3	1

Kernel (3, 3)

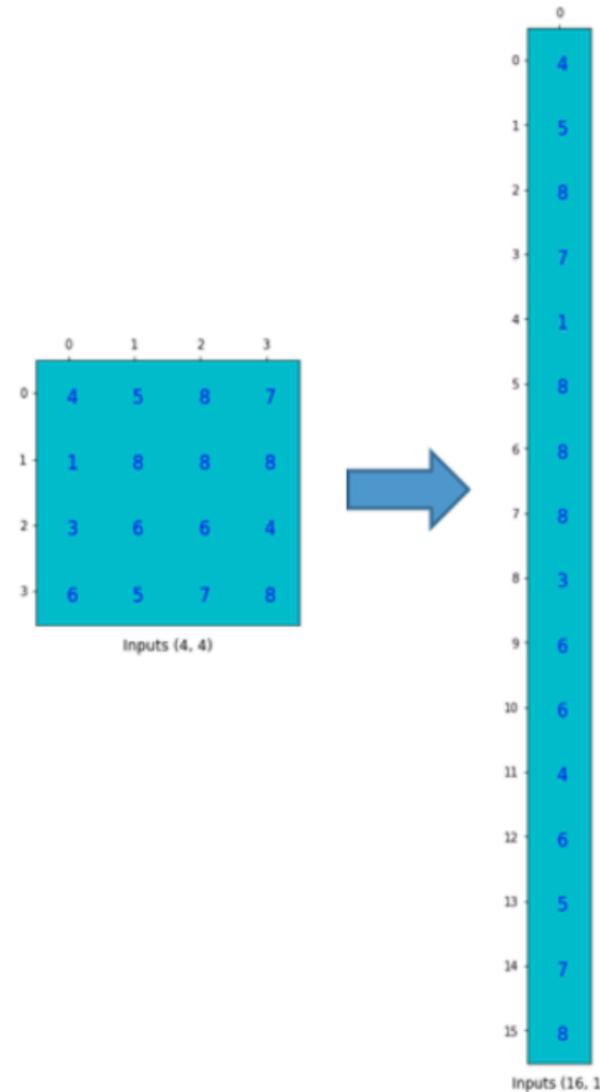
THE KERNEL CAN BE ARRANGED IN FORM OF A MATRIX:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	1	4	1	0	1	4	3	0	3	3	1	0	0	0	0	0
1	0	1	4	1	0	1	4	3	0	3	3	1	0	0	0	0
2	0	0	0	0	1	4	1	0	1	4	3	0	3	3	1	0
3	0	0	0	0	0	1	4	1	0	1	4	3	0	3	3	1

Convolution Matrix (4, 16)

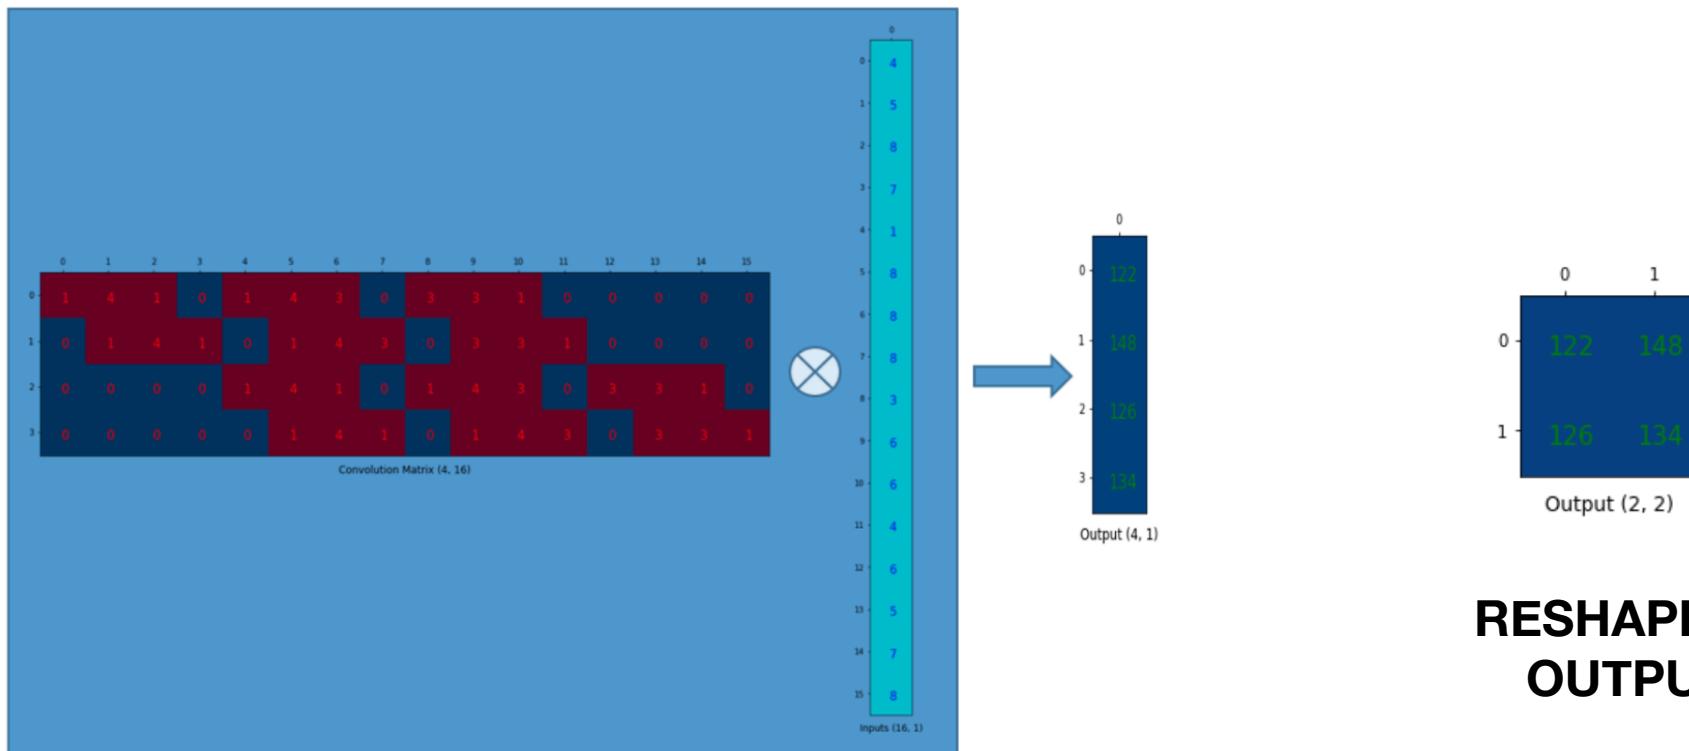
EXAMPLE TAKEN FROM HERE

## THE INPUT IS FLATTENED INTO A COLUMN VECTOR



EXAMPLE TAKEN FROM HERE

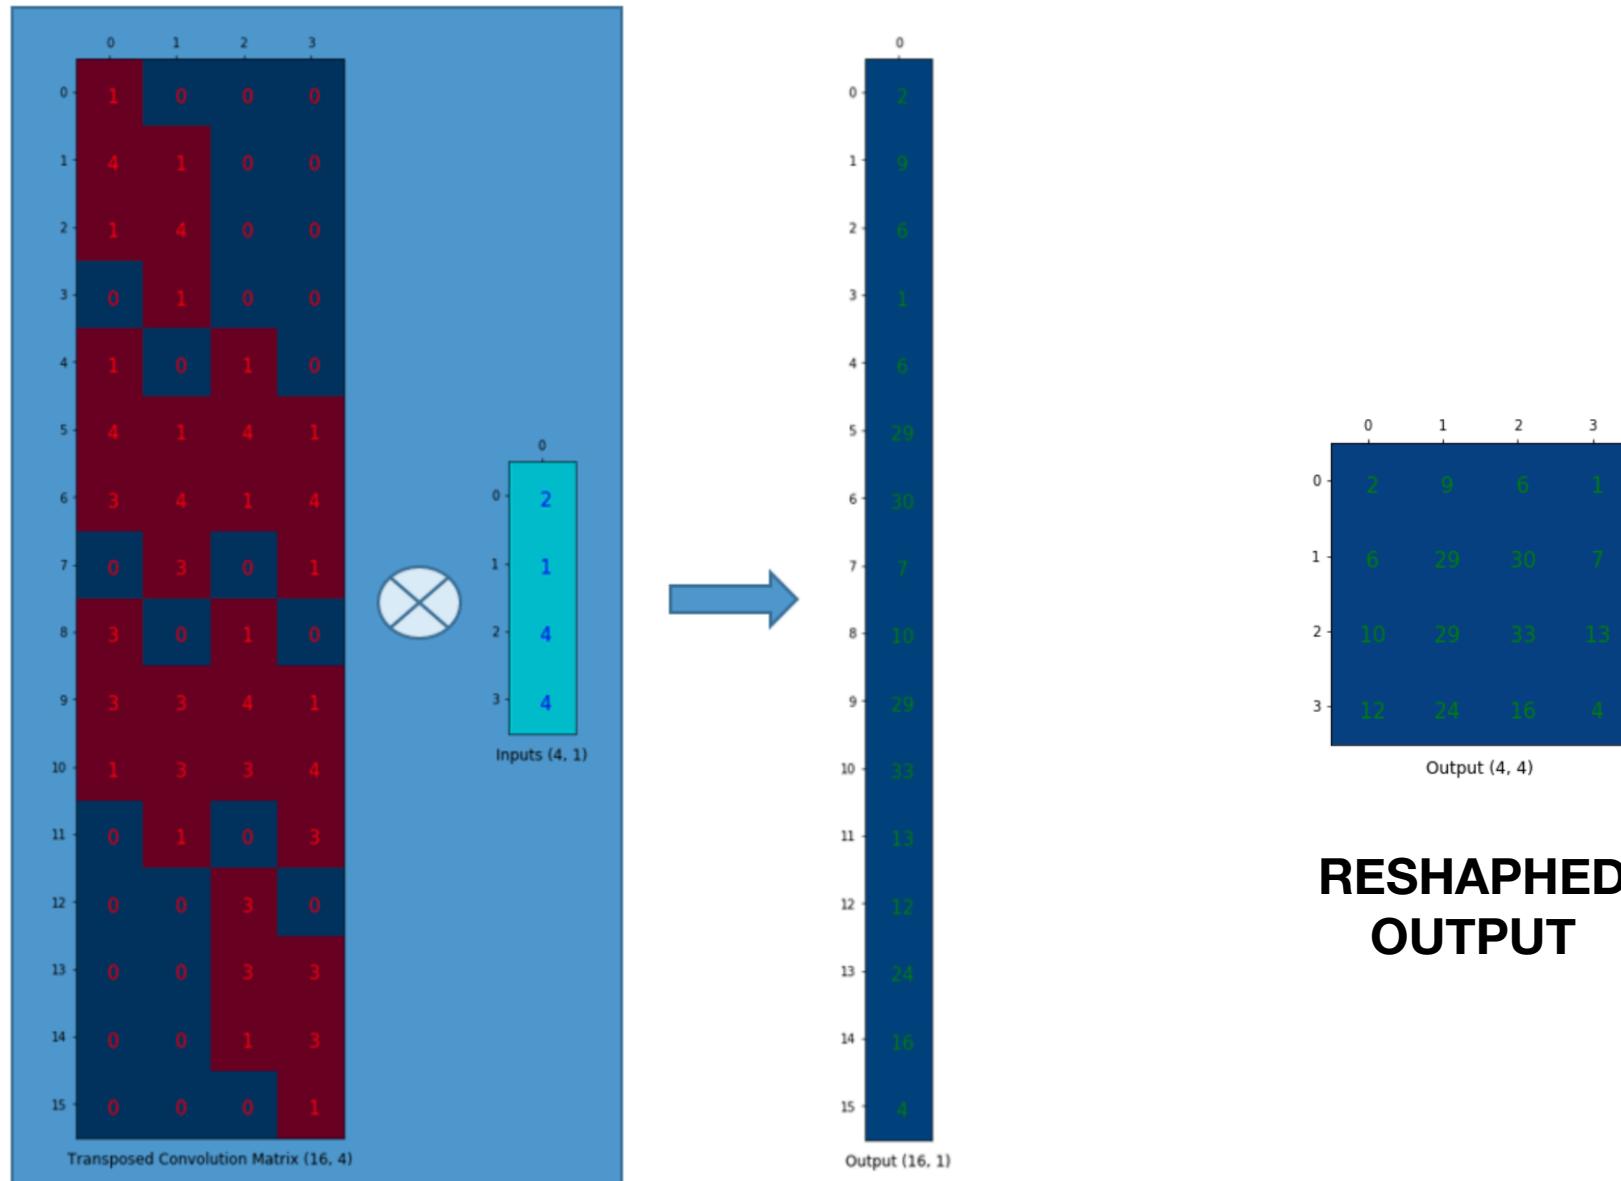
# THE CONVOLUTION IS TRANSFORMED INTO A PRODUCT OF MATRICES



**RESHAPED  
OUTPUT**

EXAMPLE TAKEN FROM HERE

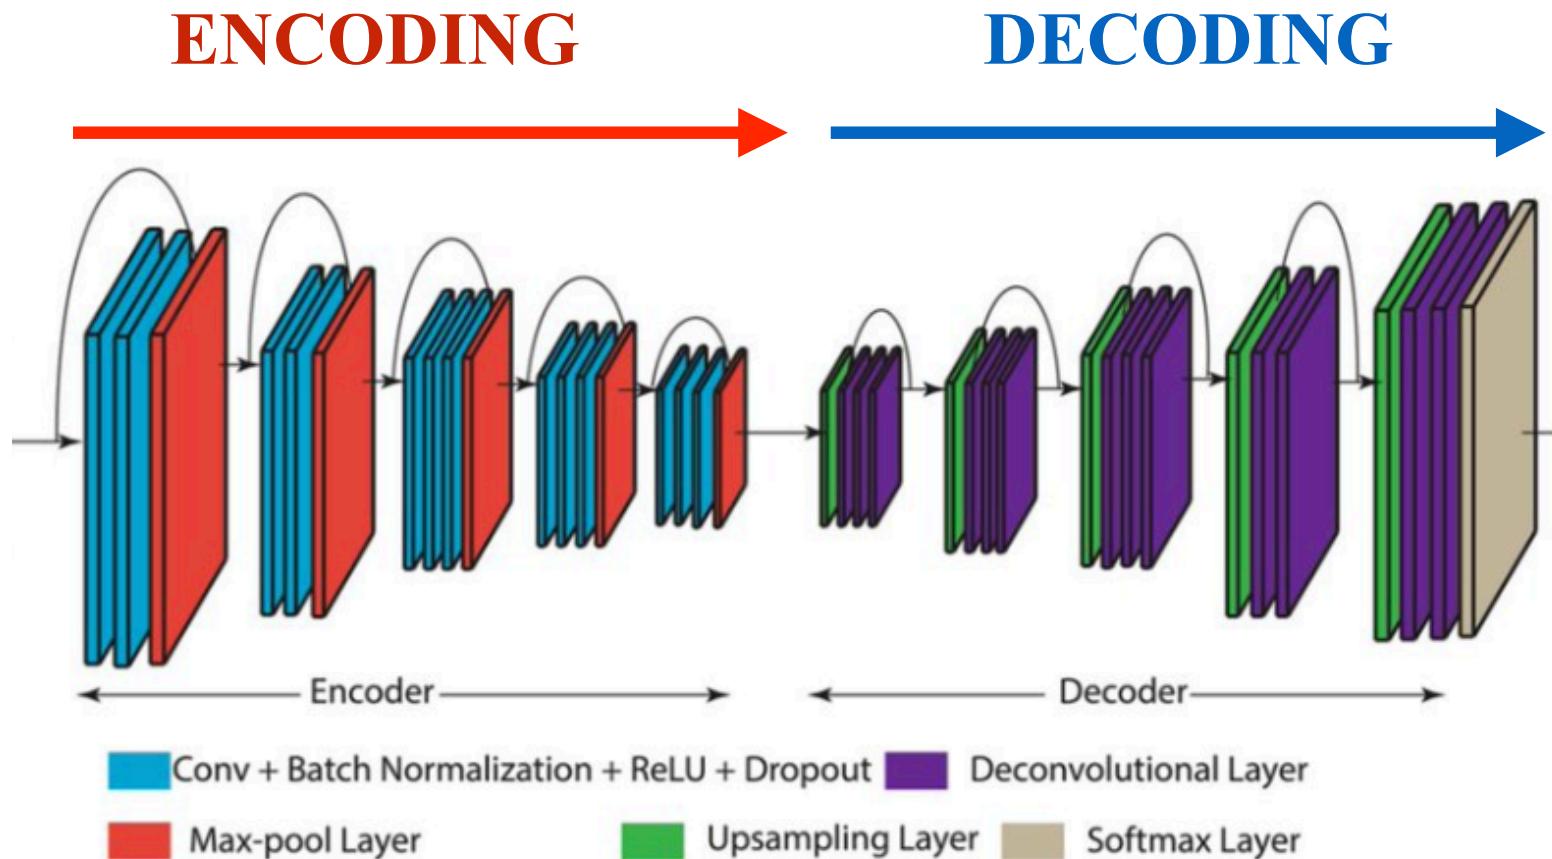
# THE TRANSPOSED CONVOLUTION IS THE INVERSE OPERATION



**RESHAPED  
OUTPUT**

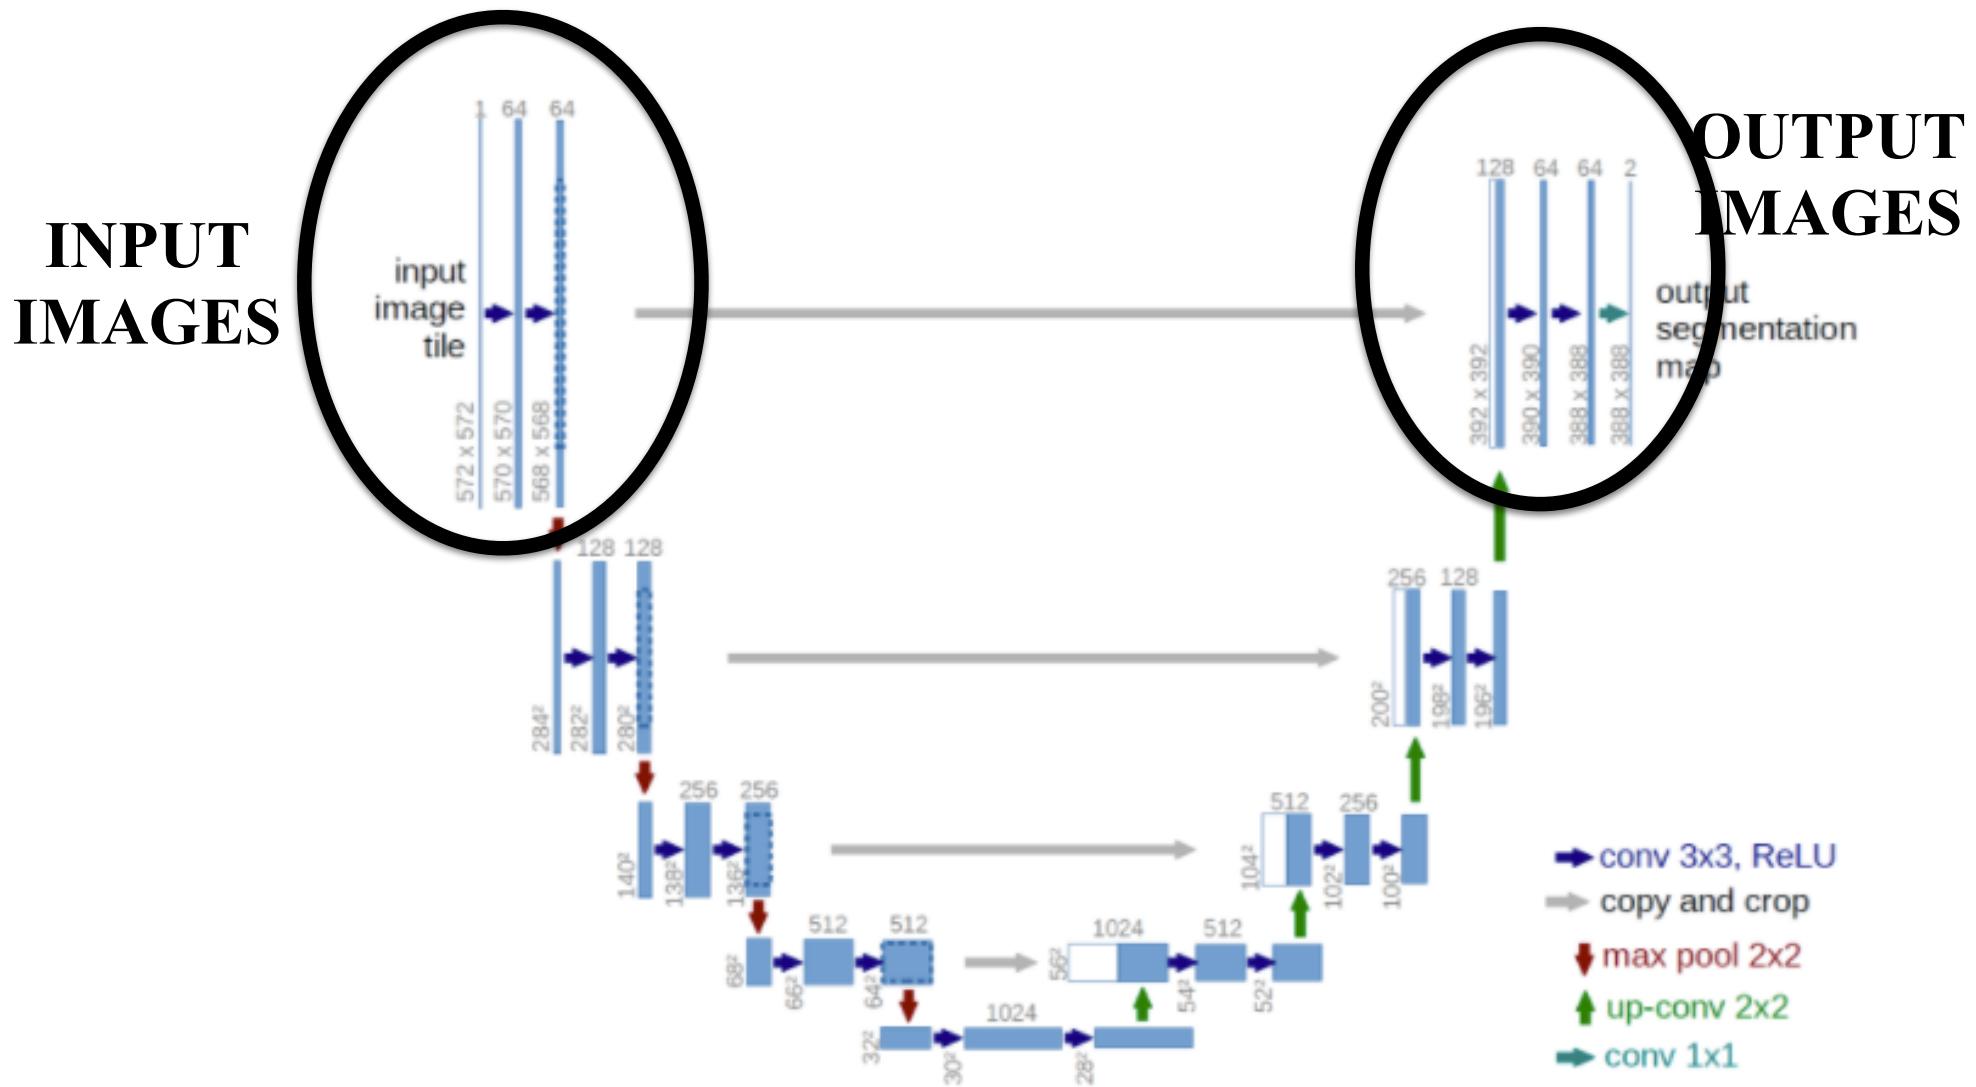
EXAMPLE TAKEN FROM HERE

# ENCODER-DECODERS GO FROM IMAGE 2 IMAGE

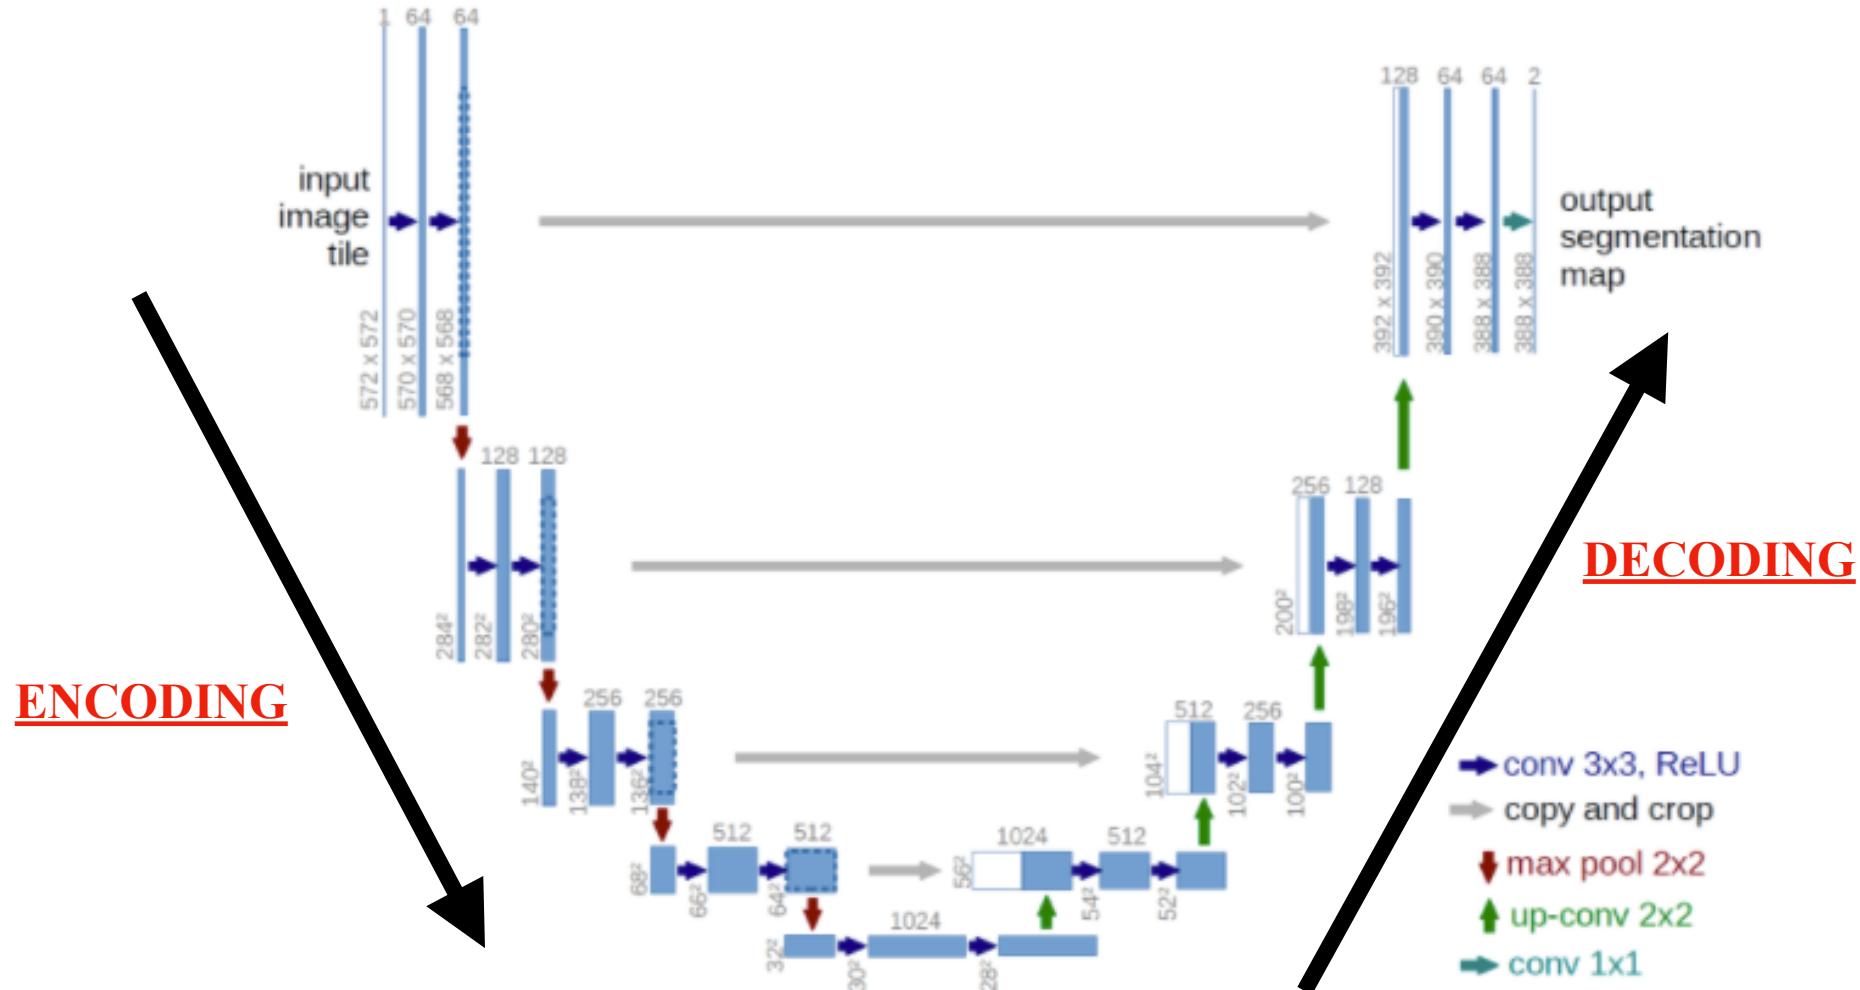


WE CALL THIS FULLY CONVOLUTIONAL  
NEURAL NETWORKS

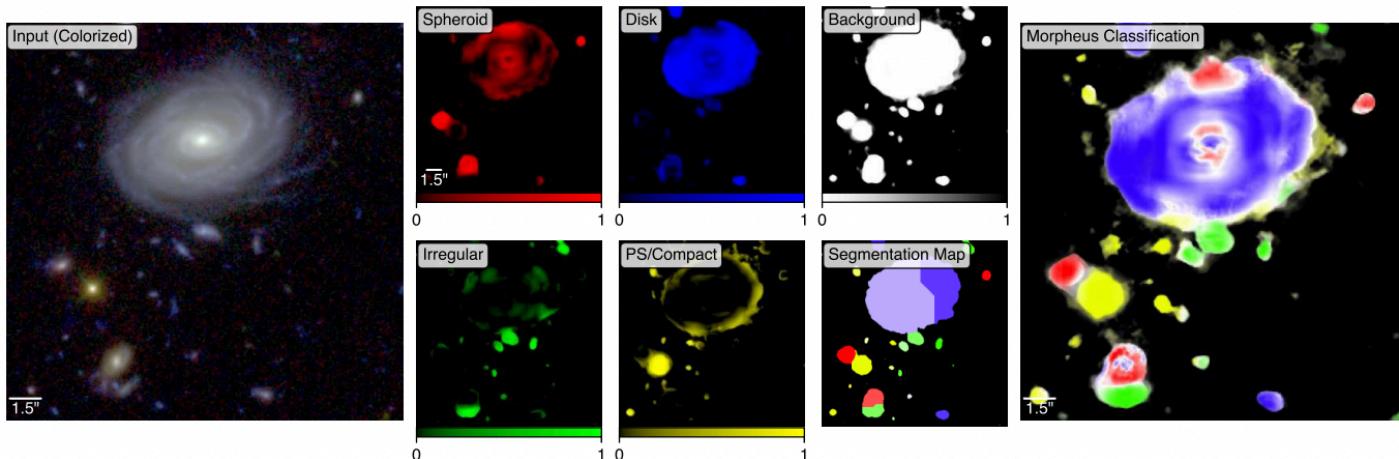
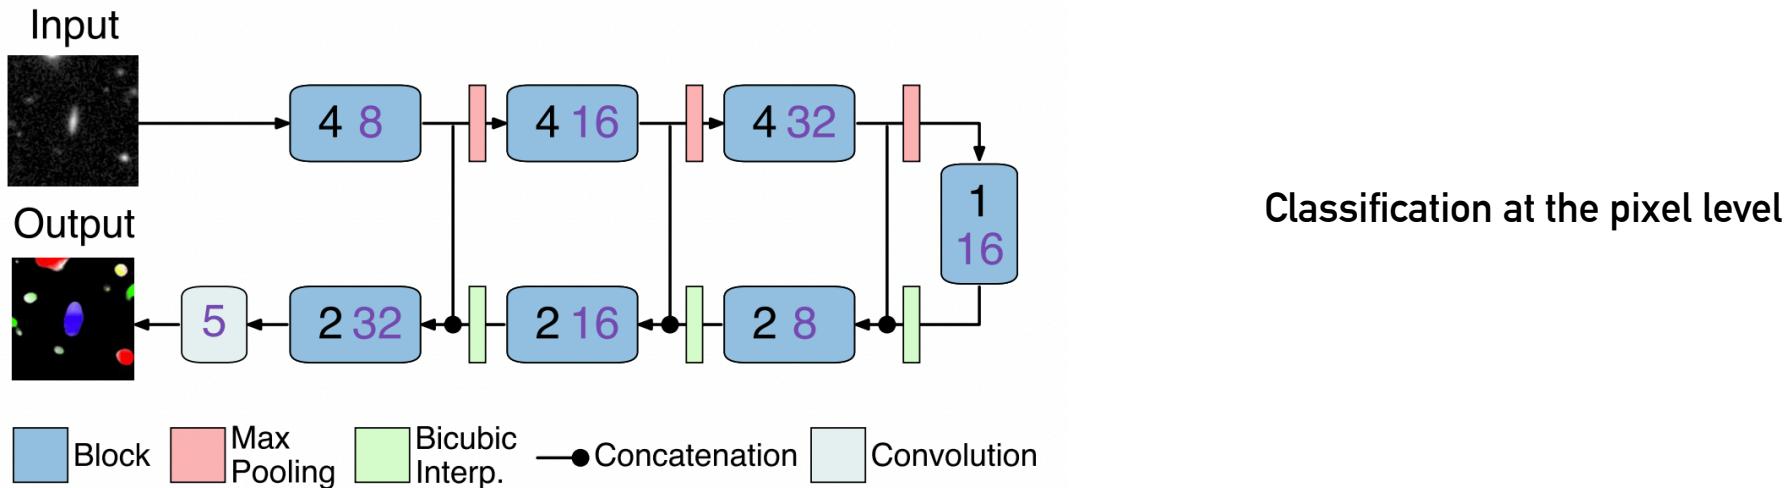
# ENCODING-DECODING TO EXTRACT IMAGE FEATURES: U-NET



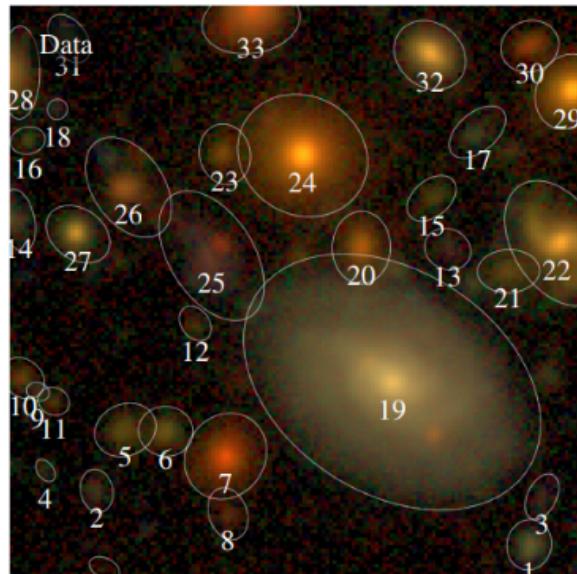
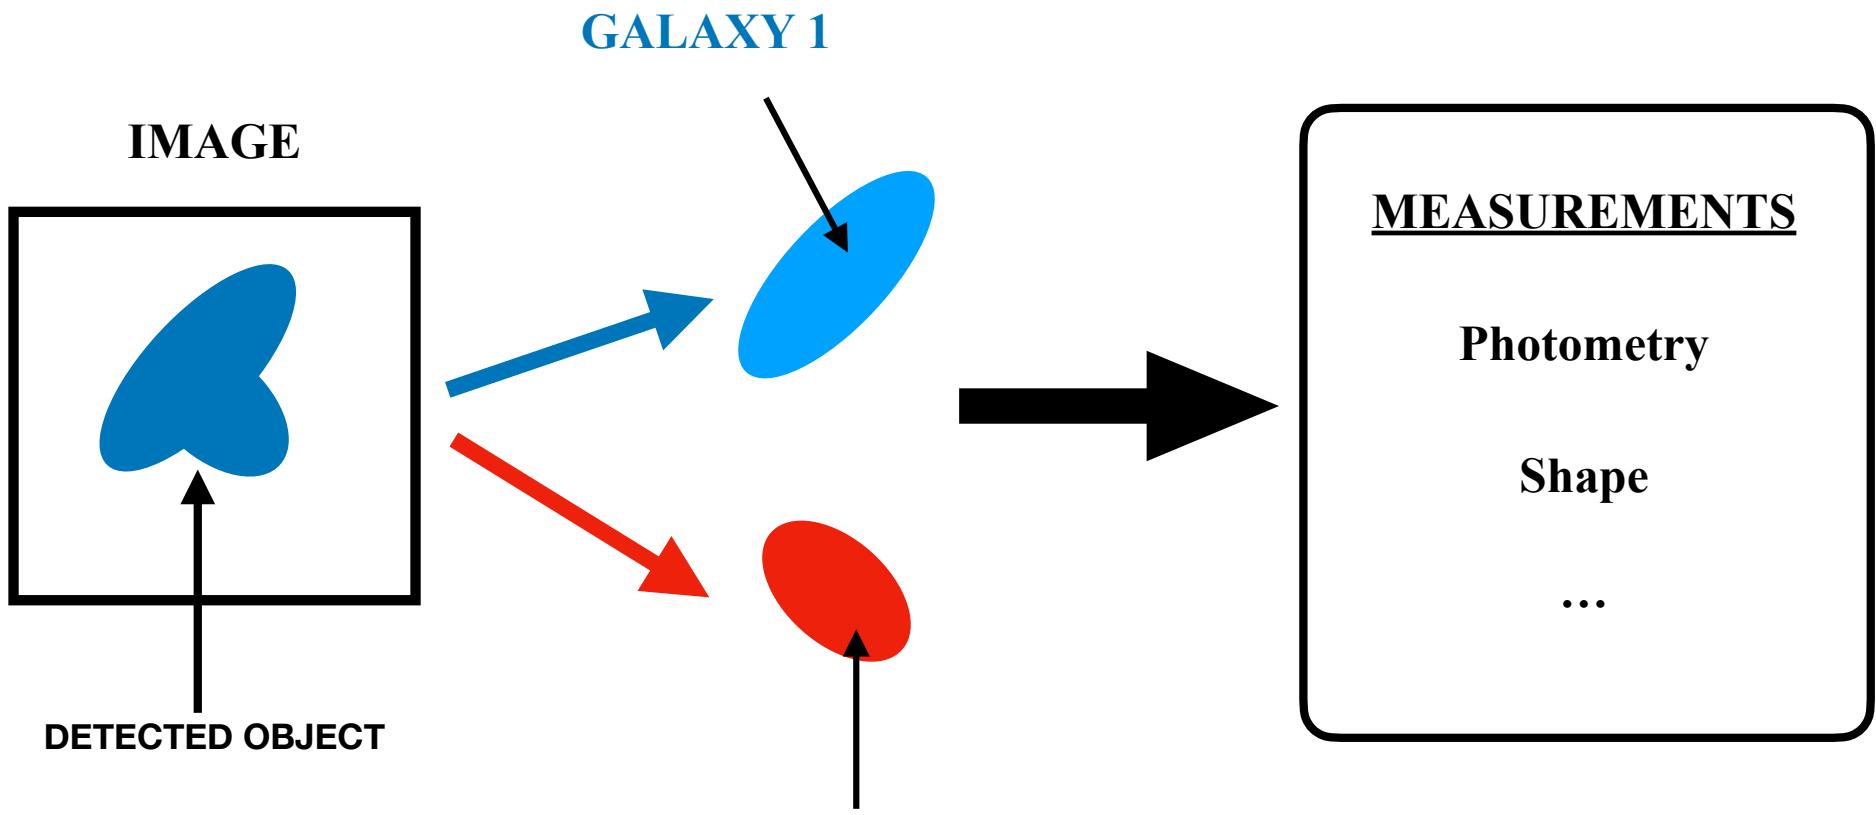
# ENCODING-DECODING TO EXTRACT IMAGE FEATURES: THE U-NET



## 2. Segmentation



Hausen+20



**>50% of objects will be affected by blending in future deep surveys such as LSST**

# U-NET (ENCODER DECODER)

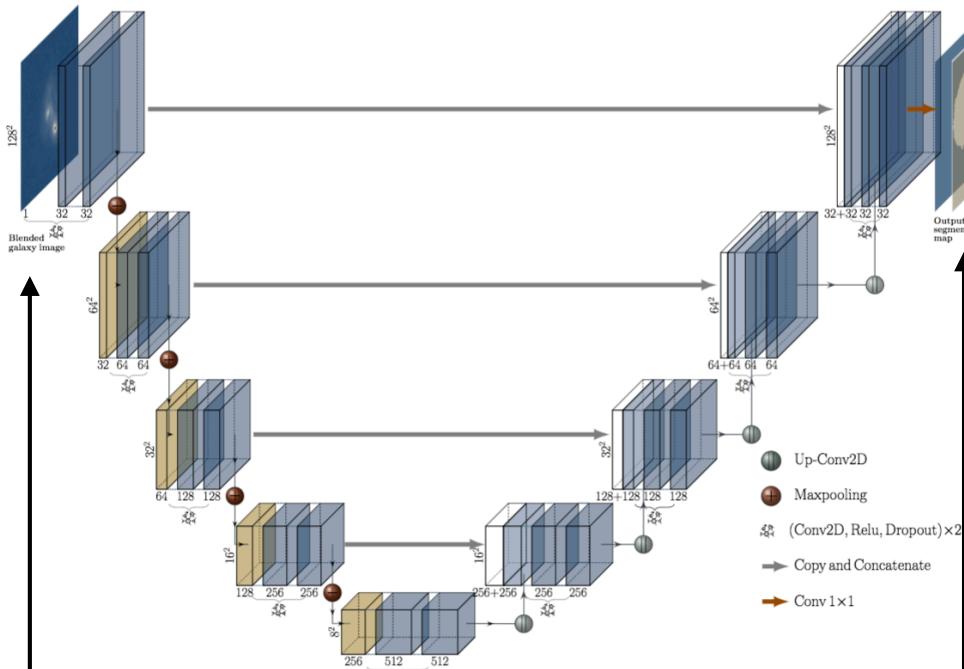
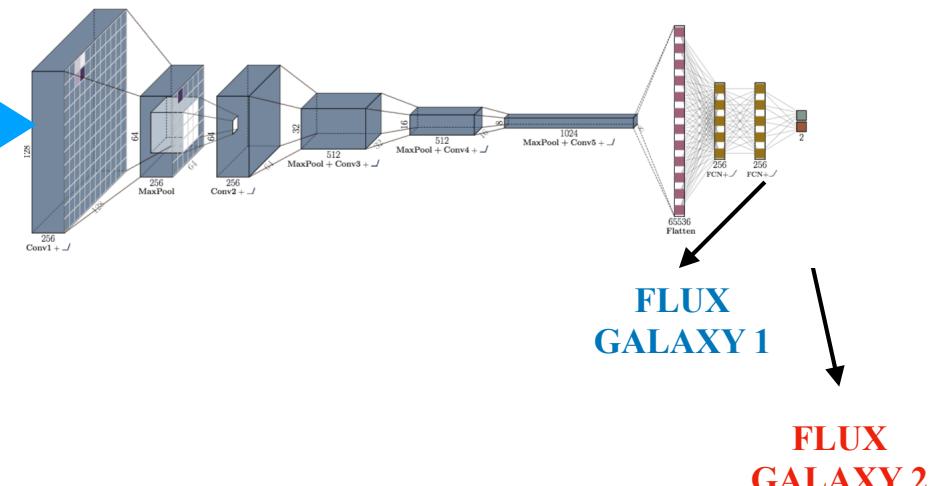
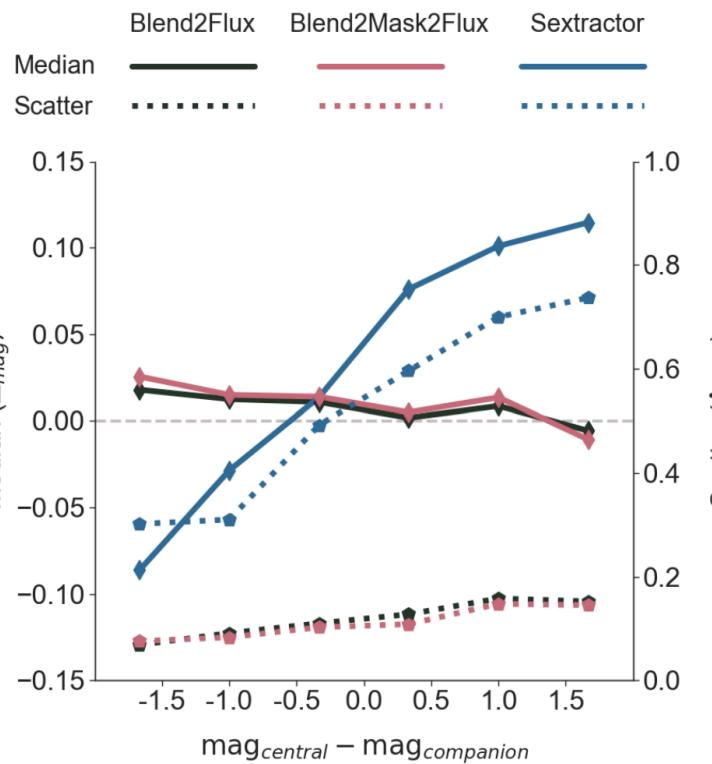
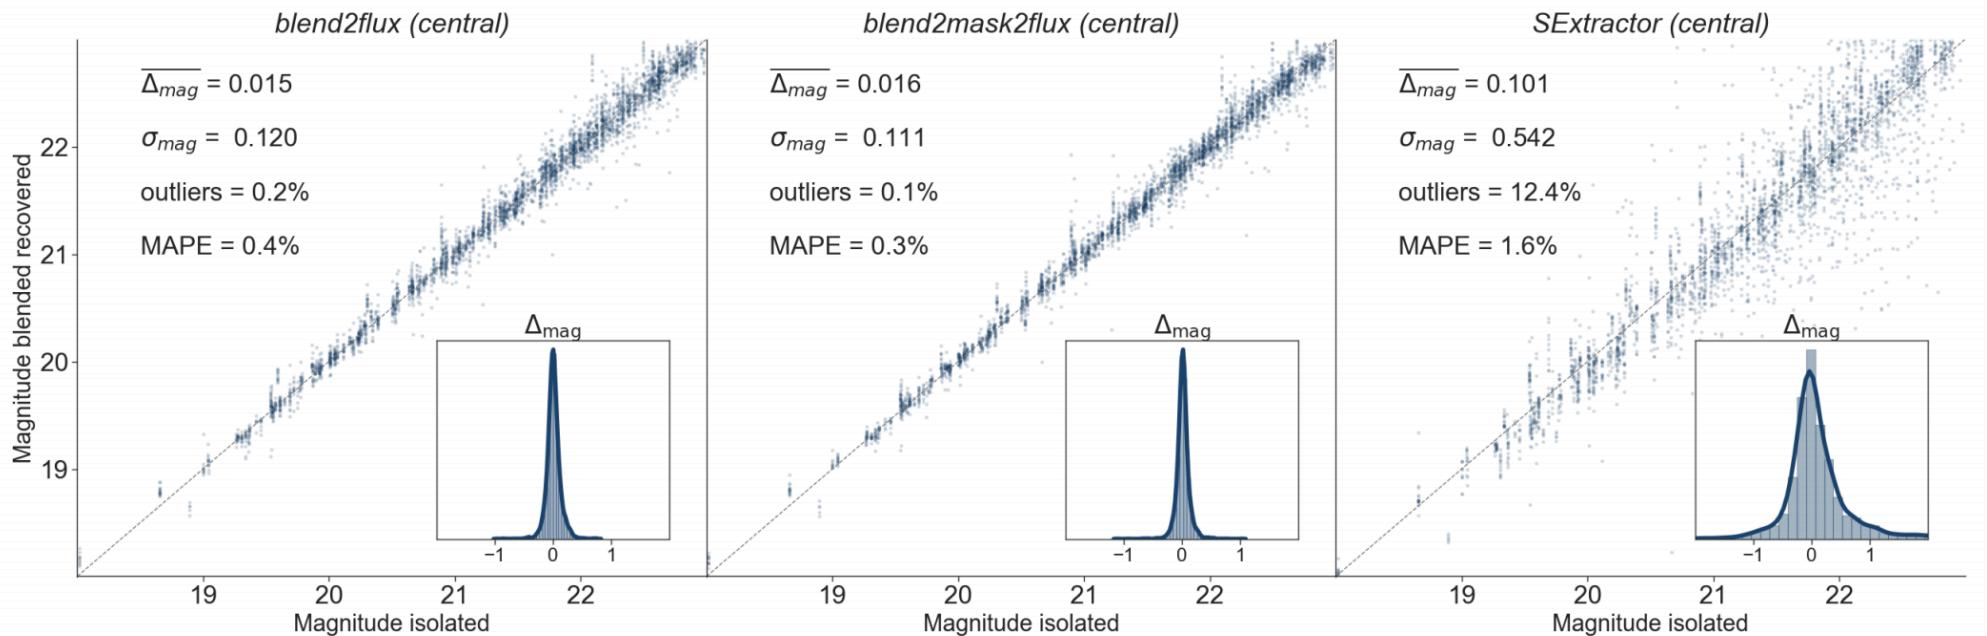


IMAGE OF  
BLENDED  
OBJECTS

OUTPUT  
SEGMENTATIO  
N MAP (BINARY  
2 CHANNELS)

# VANILLA CNN

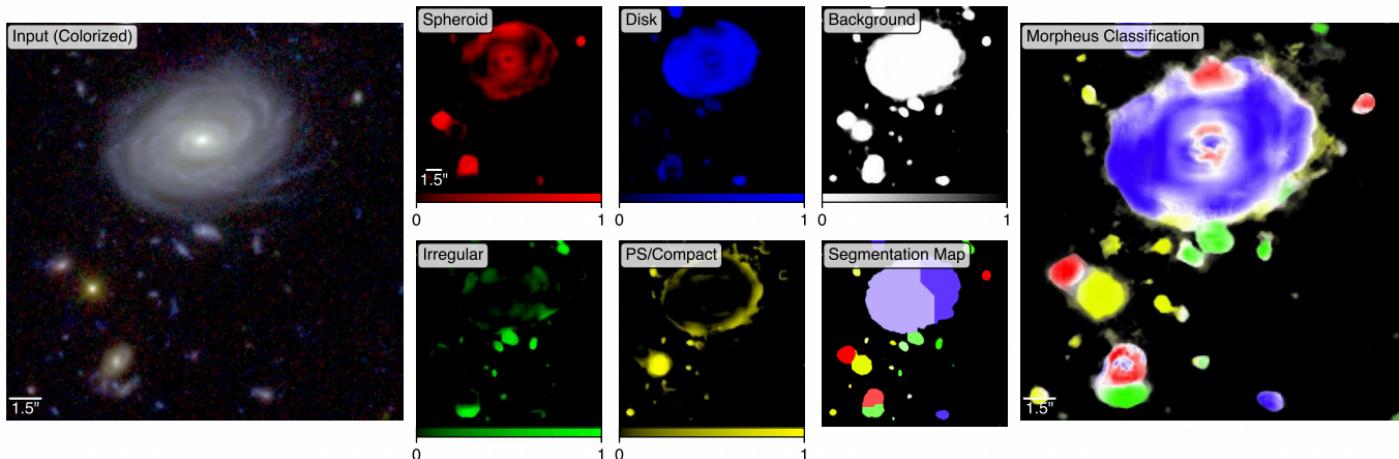
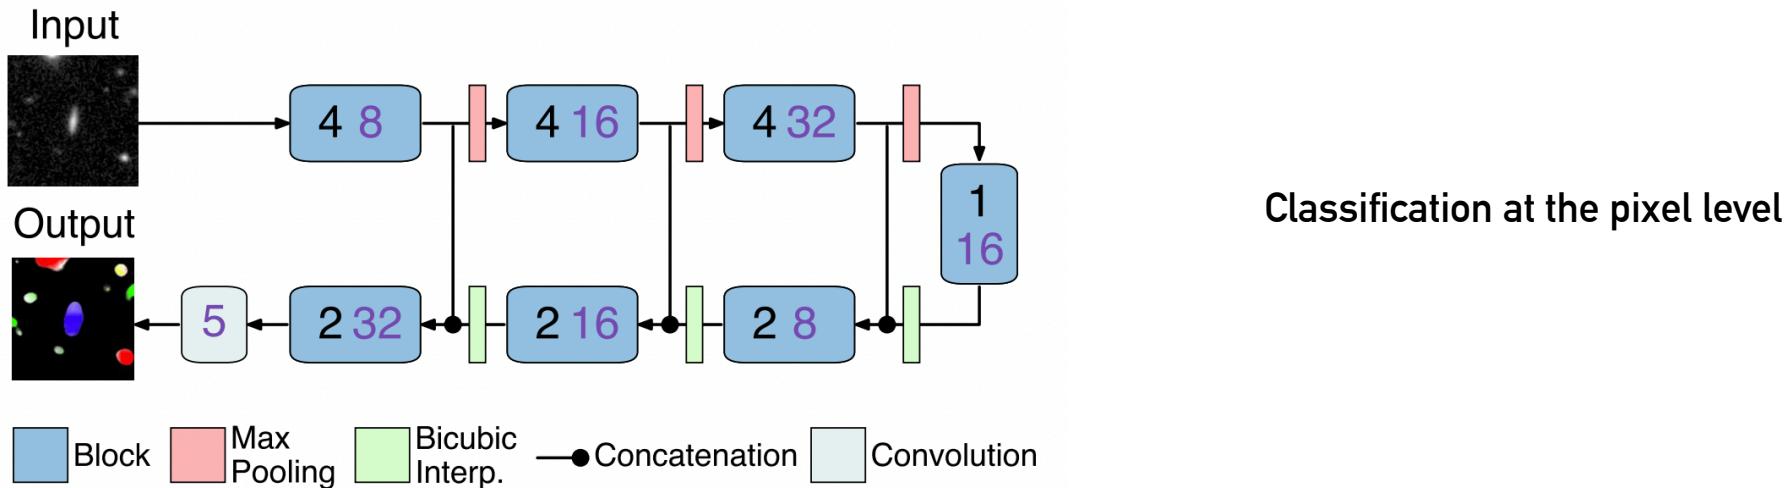




## LESSONS LEARNED:

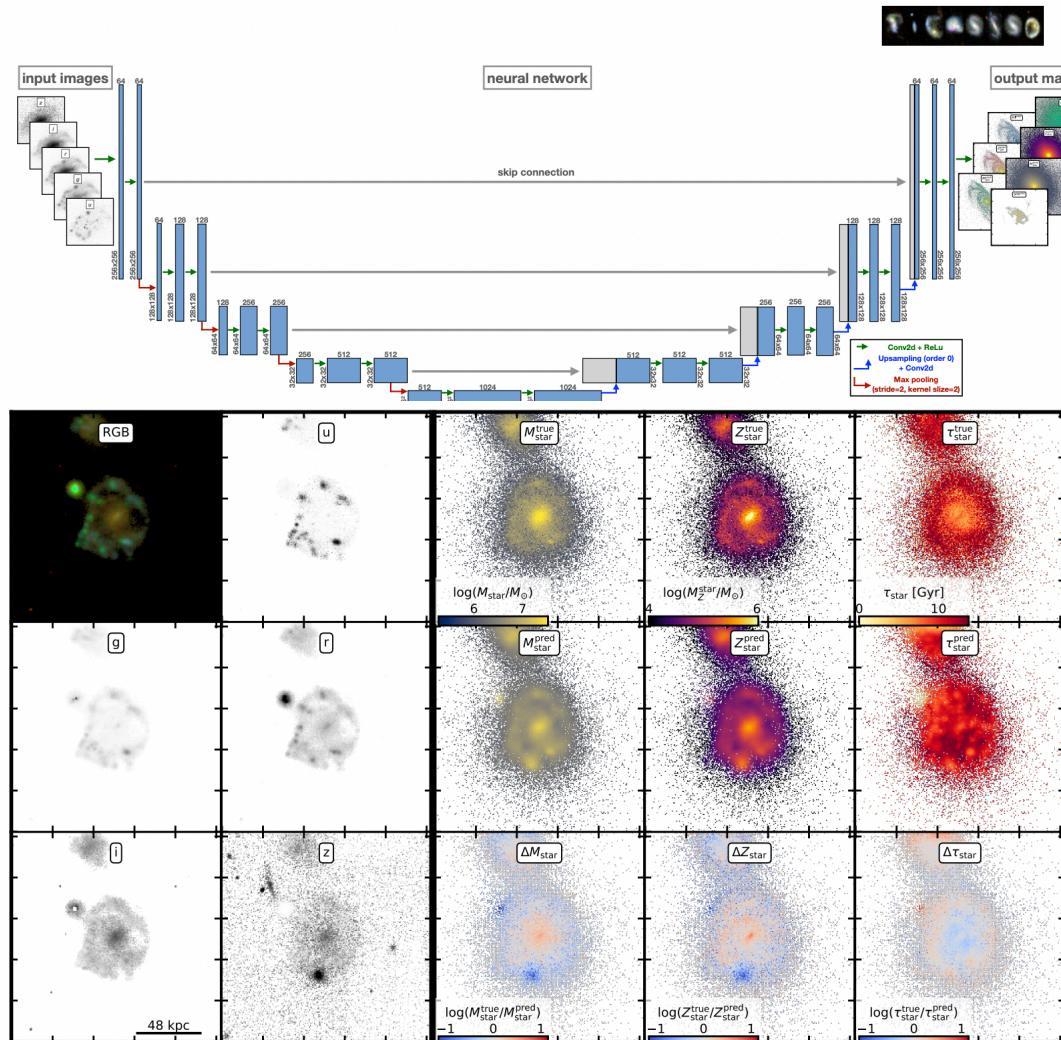
- It is easy to do better than SExtractor
- Masks do not provide additional information

## 2. Segmentation

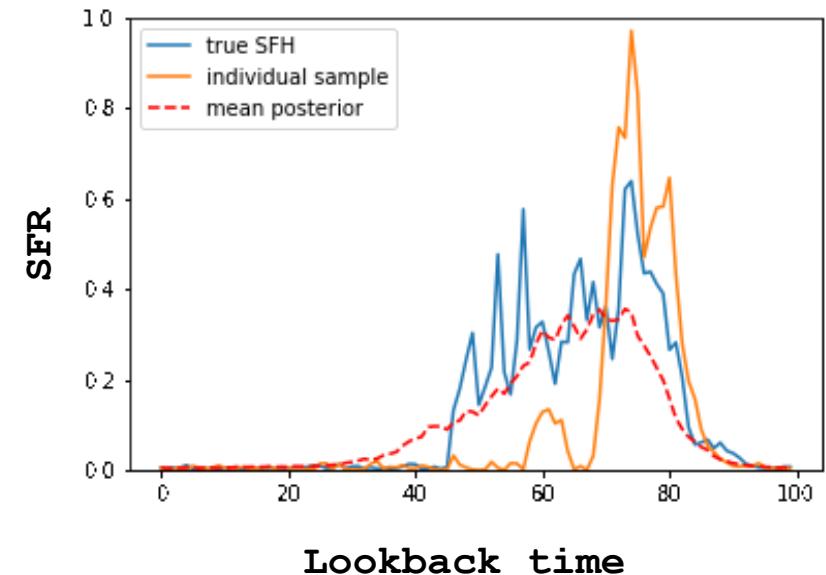
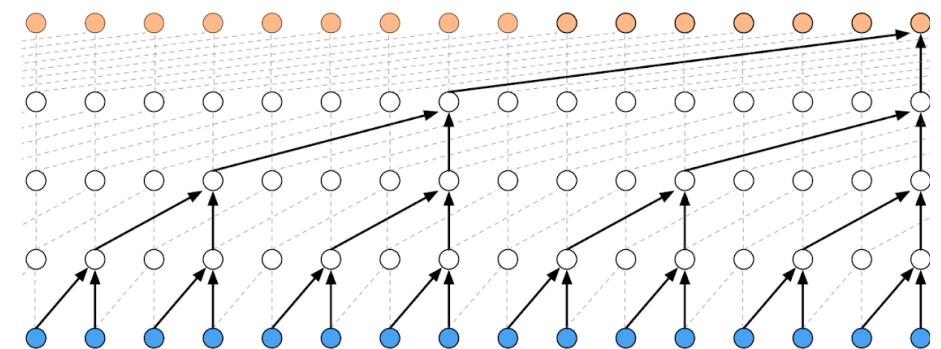


Hausen+20

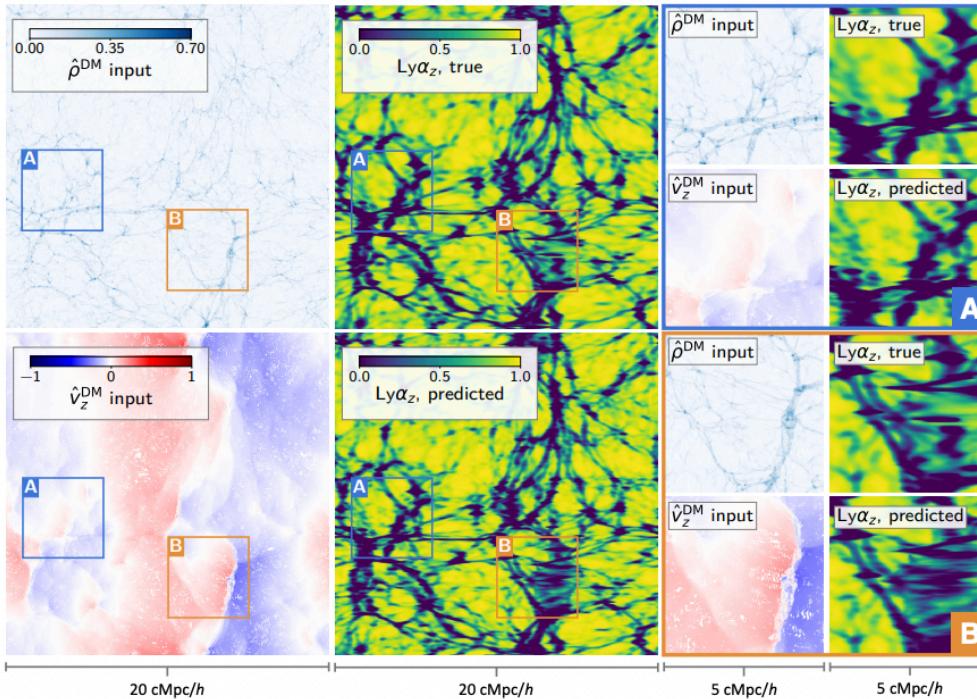
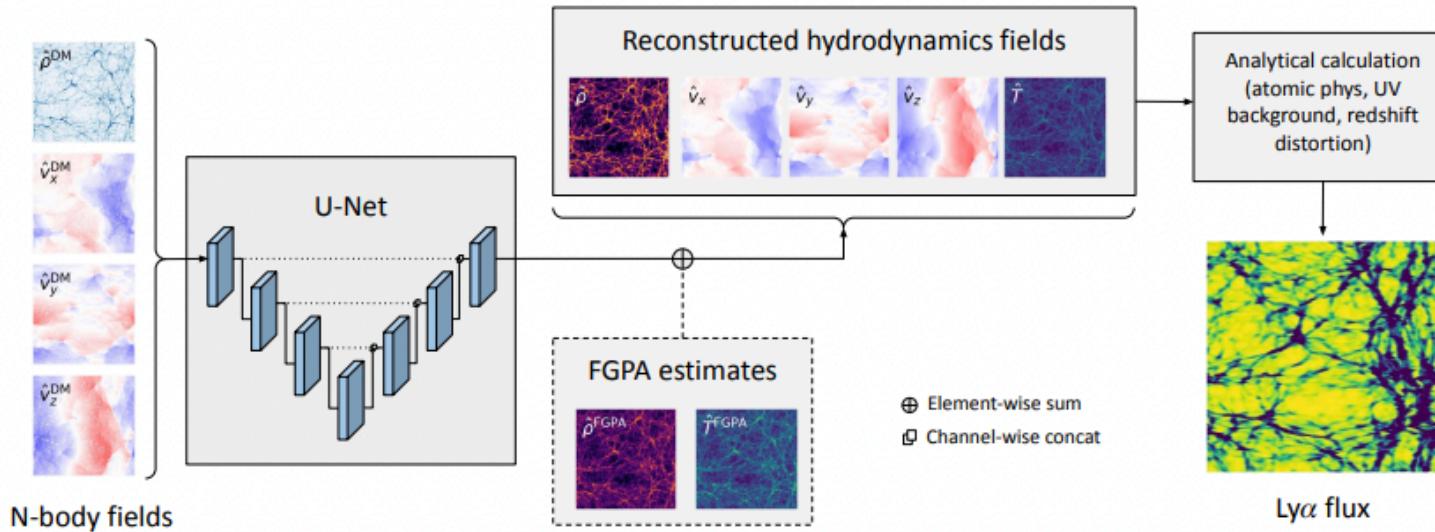
# Stellar Populations



**Buck+21**



# Painting Baryons



**Harrington+21**

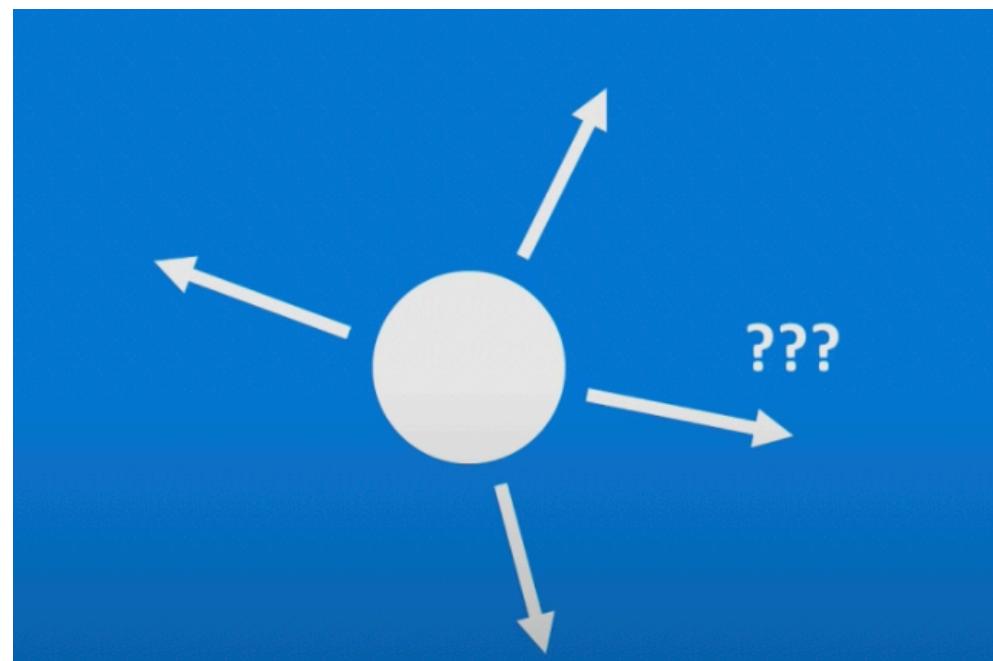
**Neural Networks are used to learn the non-linear mapping between cheap dark matter only simulations to expensive baryonic physics**

Rodriguez+19, Modi+18, Berger+18, He+18, Zhang+19, Troster+19, Zamudio-Fernandez+19, Perraudin+19, Charnock+19, List+19, Giusarma+19, Bernardini+19, Chardin+19, Mustafa+19, Ramanah+20, Tamasiunas+20, Feder+20, Moster+20, Thiele+20, Wadekar+20, Dai+20, Li+20, Lucie-Smith+20, Kasmanoff+20, Ni+21, Rouhaiainen+21, Harrington+21, Horowitz+21, Horowitz+21, Bernardini+21, Schaurecker+21, Etezad-Razavi+21, Curtis+21

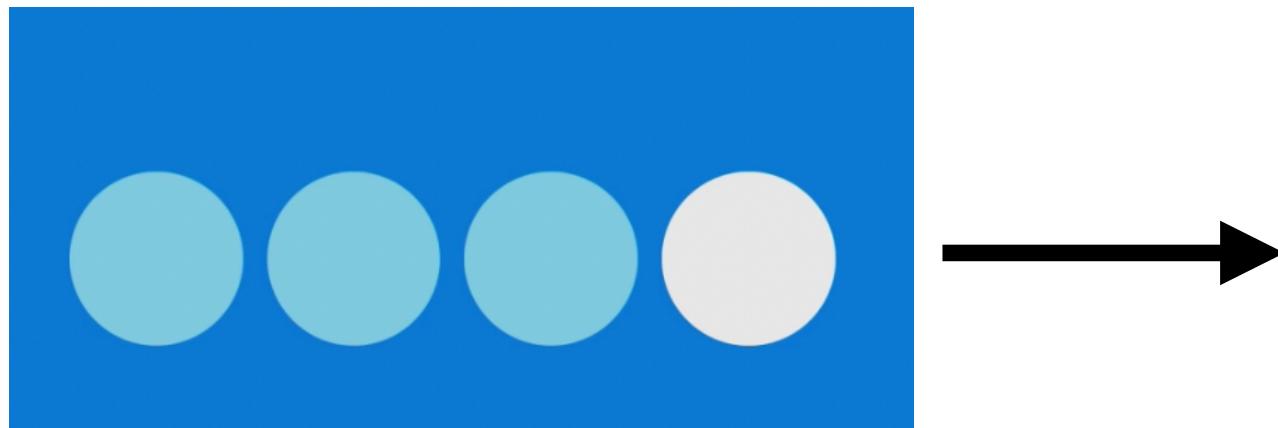
# **SEQUENCE MODELLING**

\*based on MIT Lecture by Ava Soleimany

GIVEN AN IMAGE OF A BALL, CAN YOU PREDICT  
WHERE IT WILL GO NEXT?



GIVEN AN IMAGE OF A BALL, CAN YOU PREDICT  
WHERE IT WILL GO NEXT?



Previous positions help guessing the future

# LET'S TAKE A SIMPLE EXAMPLE OF LANGUAGE MODELLING

“This morning I took my cat for a walk.”

given these words

predict the  
next word

# LET'S TAKE A SIMPLE EXAMPLE OF LANGUAGE MODELLING

“This morning I took my cat for a walk.”

given these words

predict the  
next word

Normal ANNs and CNNs cannot handle variable length inputs ...

# WE COULD SIMPLY USE A FIXED WINDOW...

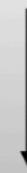
“This morning I took my cat for a walk.”

given these words

predict the  
next word

[ 1 0 0 0 0 0 1 0 0 0 ]

for            a

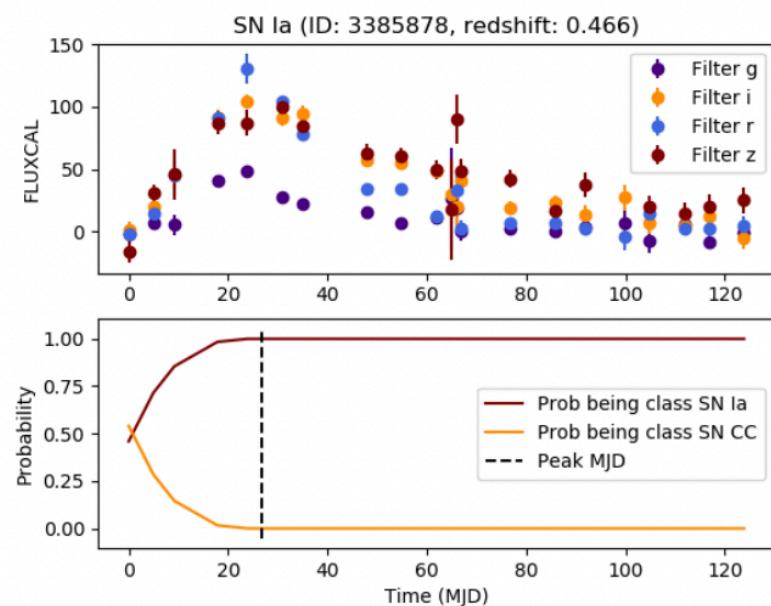


prediction

# HOWEVER THIS CANNOT HANDLE LONG TERM MEMORY

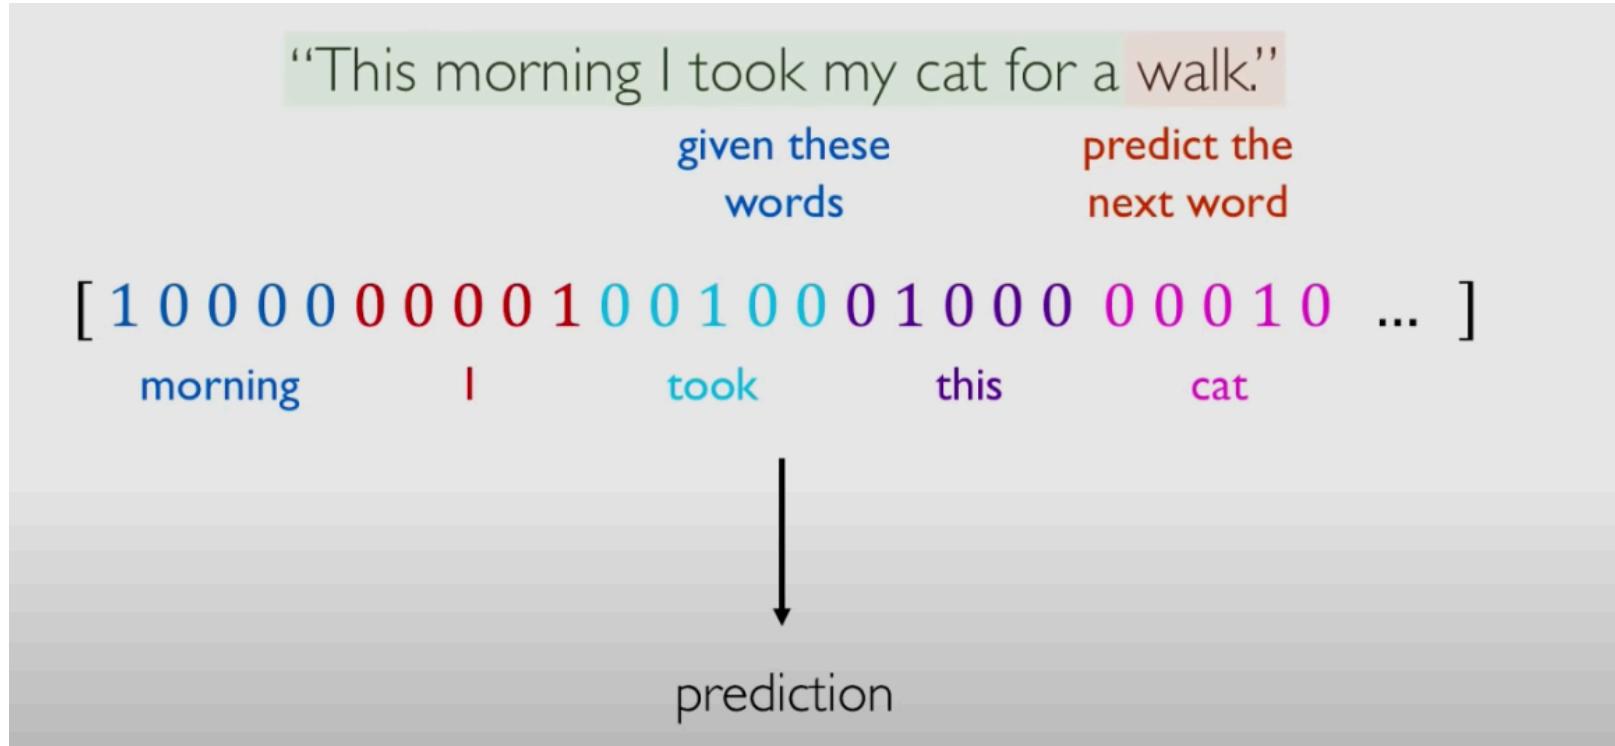
“France is where I grew up, but I now live in Boston. I speak fluent \_\_\_\_.”

IN SOME CASES, INFORMATION FROM THE DISTANT PAST  
IS NEEDED FOR A CORRECT PREDICTION



Moller+19

# ANOTHER ALTERNATIVE COULD BE TO FEED A REALLY LARGE WINDOW



BUT WE STILL HAVE THE PROBLEM THAT THERE IS NO PARAMETER SHARING (SAME AS PIXELS)

# **RECURRENT NEURAL NETWORKS (RNNs)**

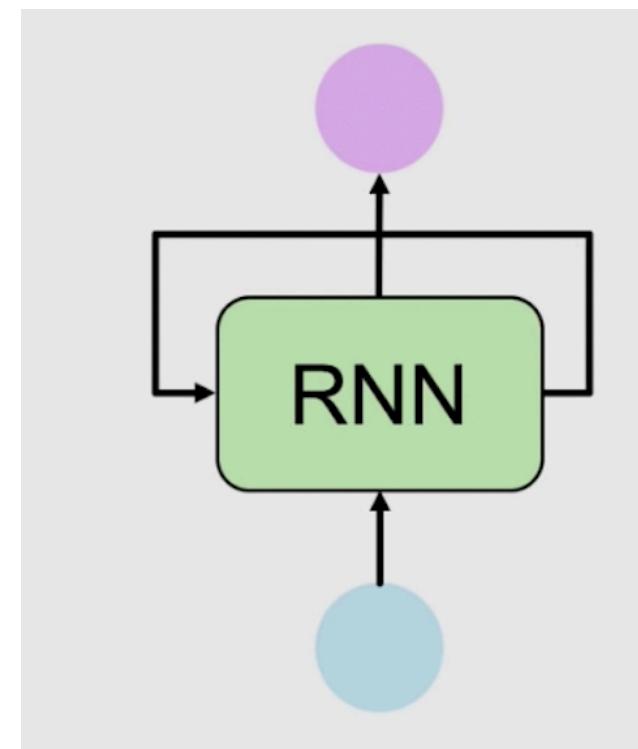
## OUR WISH LIST:

1. Handle variable-length sequences
2. Track long term dependencies
3. Maintain information about order
4. Share parameters across the sequence

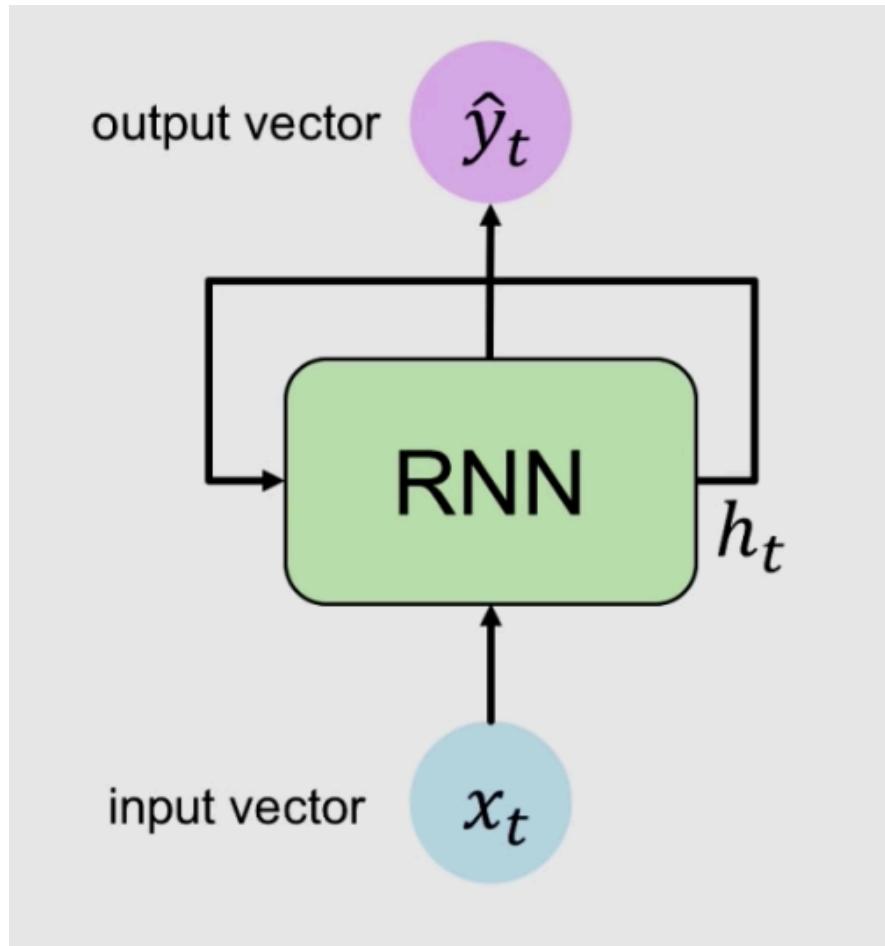
## OUR WISH LIST:

1. Handle variable-length sequences
2. Track long term dependencies
3. Maintain information about order
4. Share parameters across the sequence

Recurrent Neural Networks  
offer a first solution to this problem



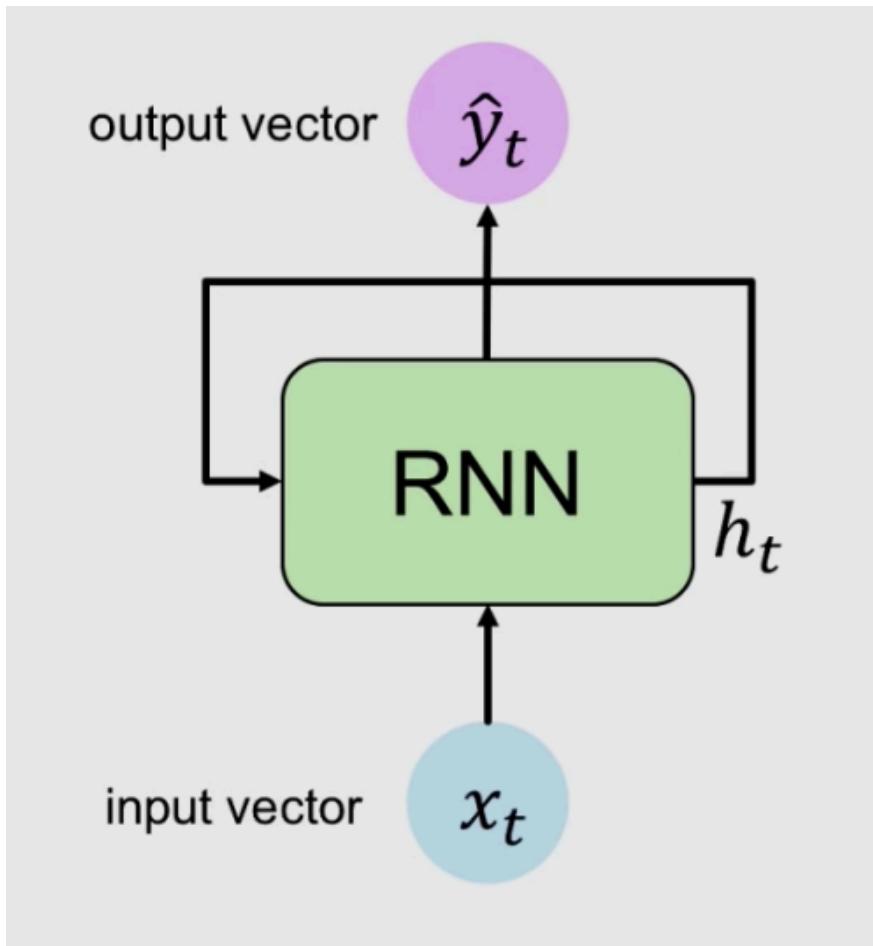
# The RNN Block



$$h_t = f_W(h_{t-1}, x_t)$$

cell state      function parameterized by W      old state      input vector at time step t

# The RNN Block



Output Vector

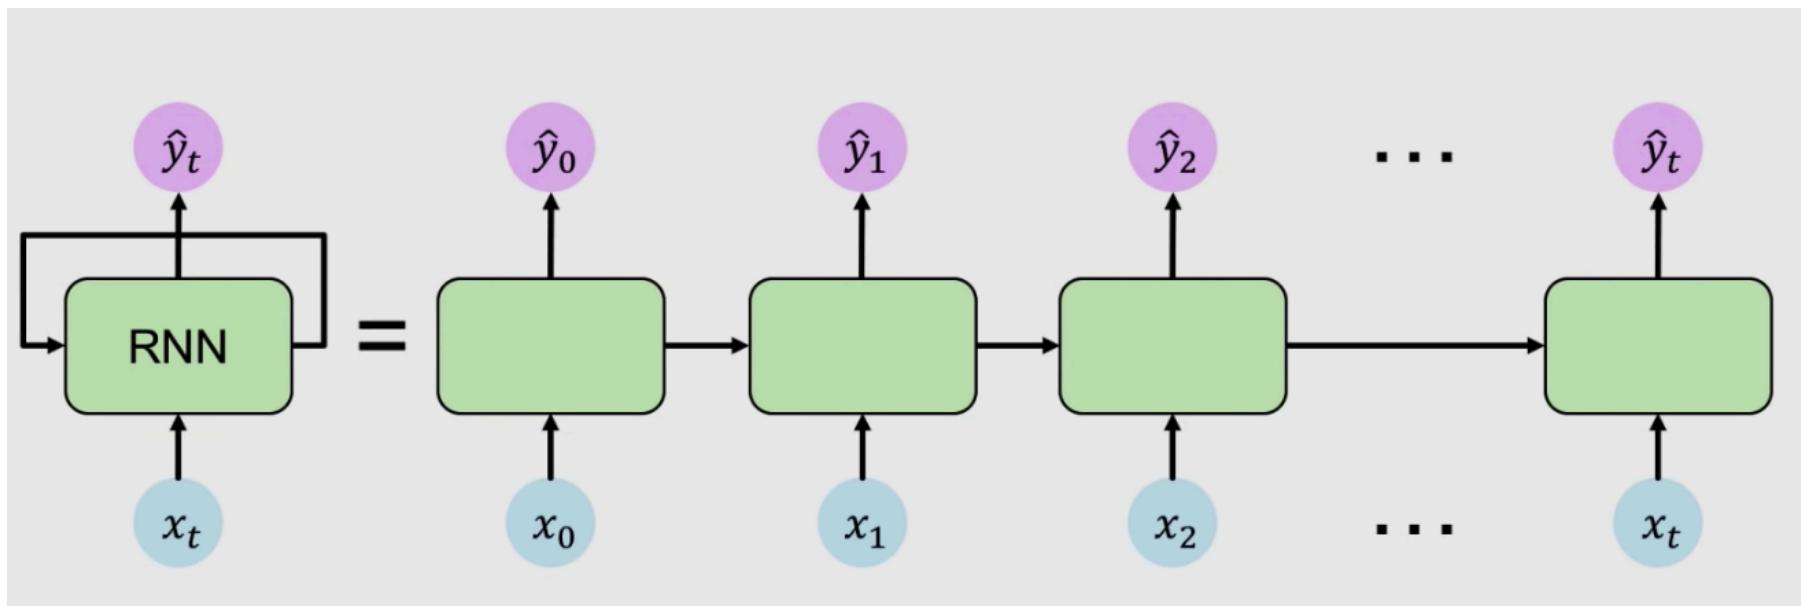
$$\hat{y}_t = \mathbf{W}_{hy}^T h_t$$

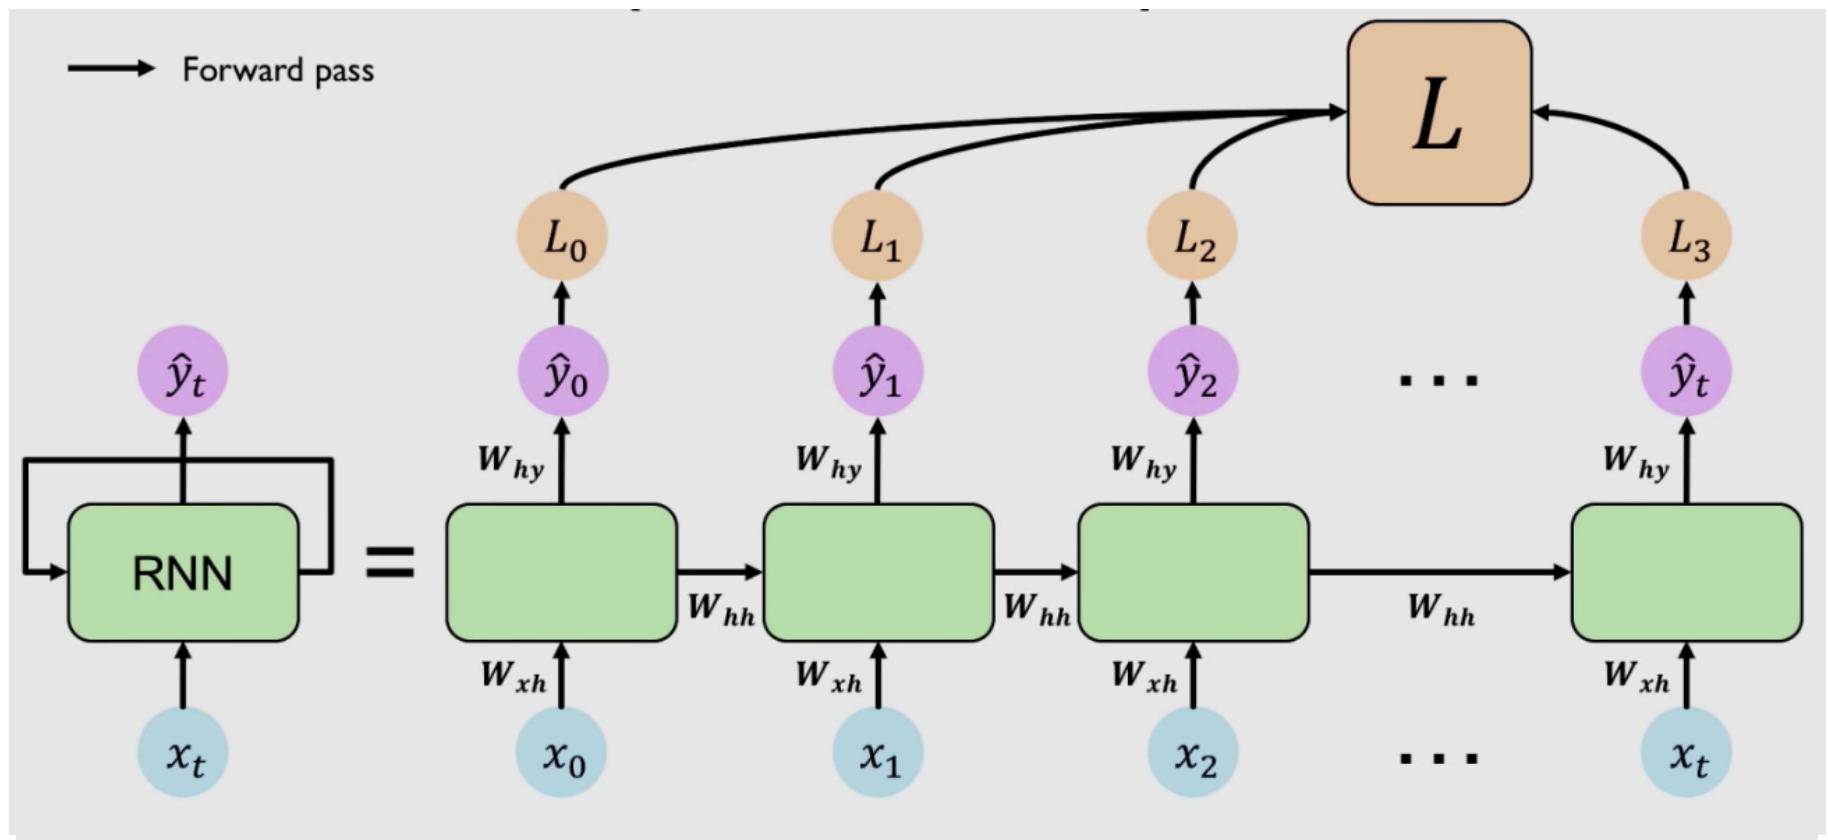
Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}^T h_{t-1} + \mathbf{W}_{xh}^T x_t)$$

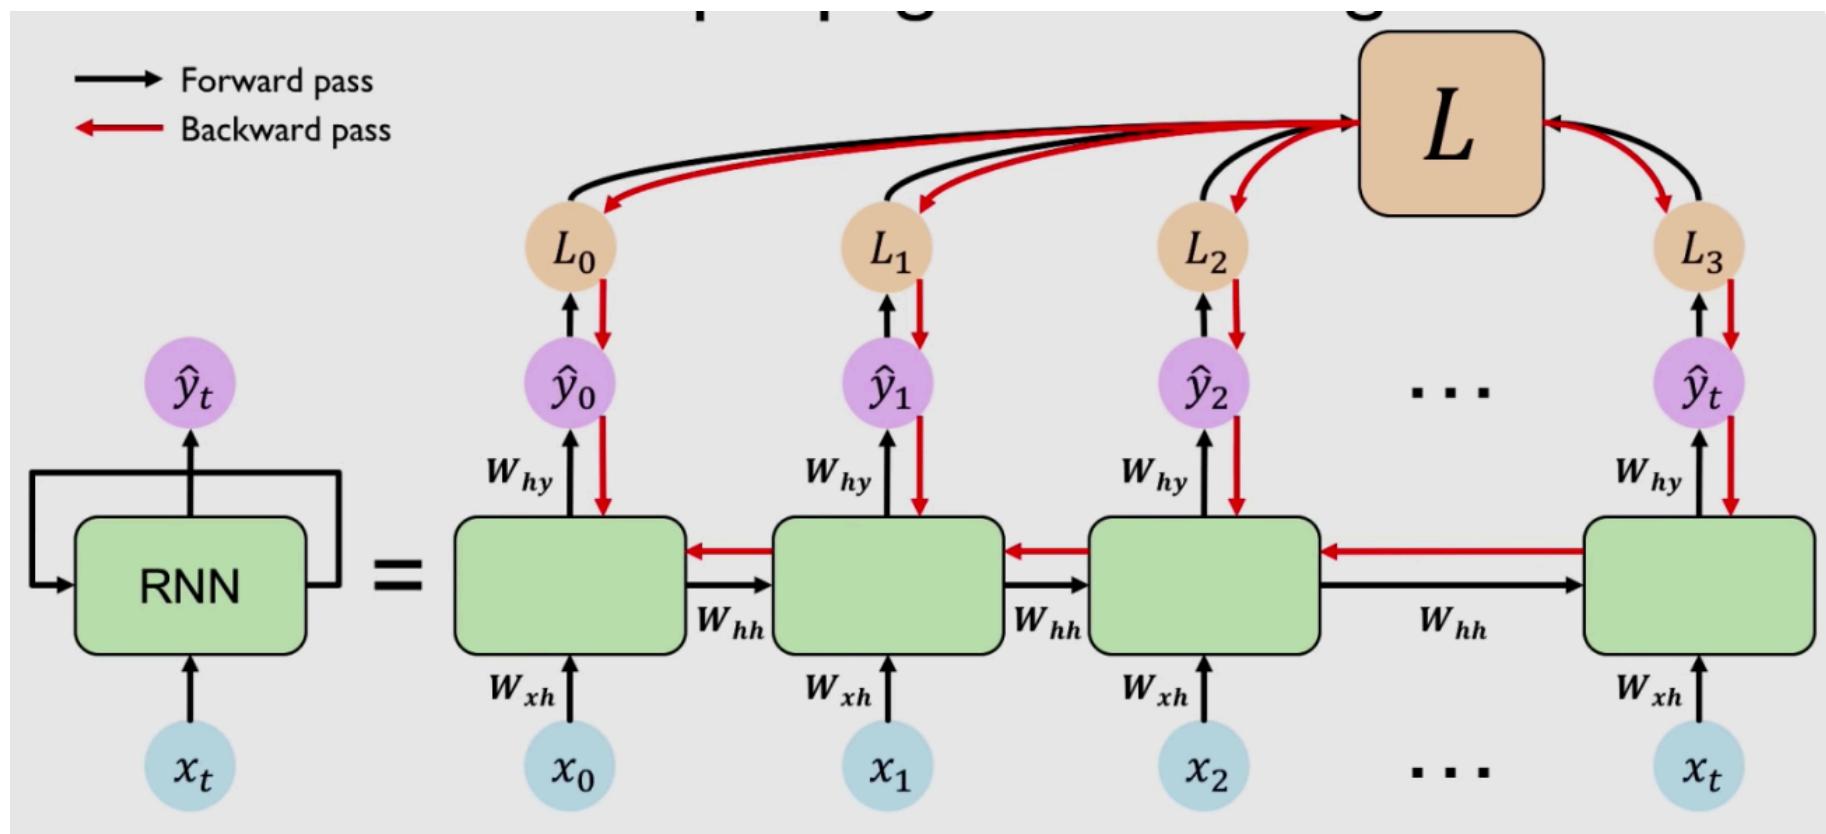
Input Vector

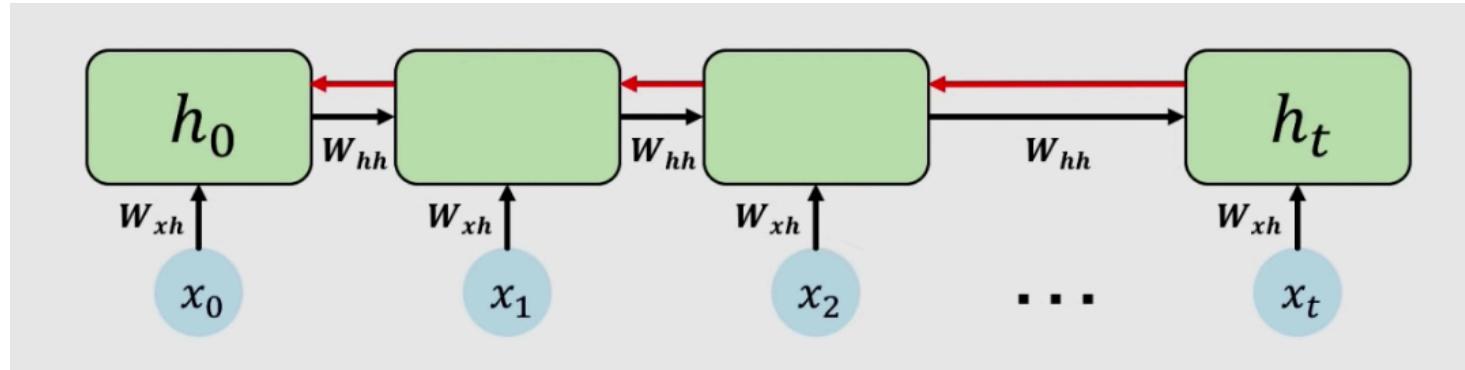
$$x_t$$



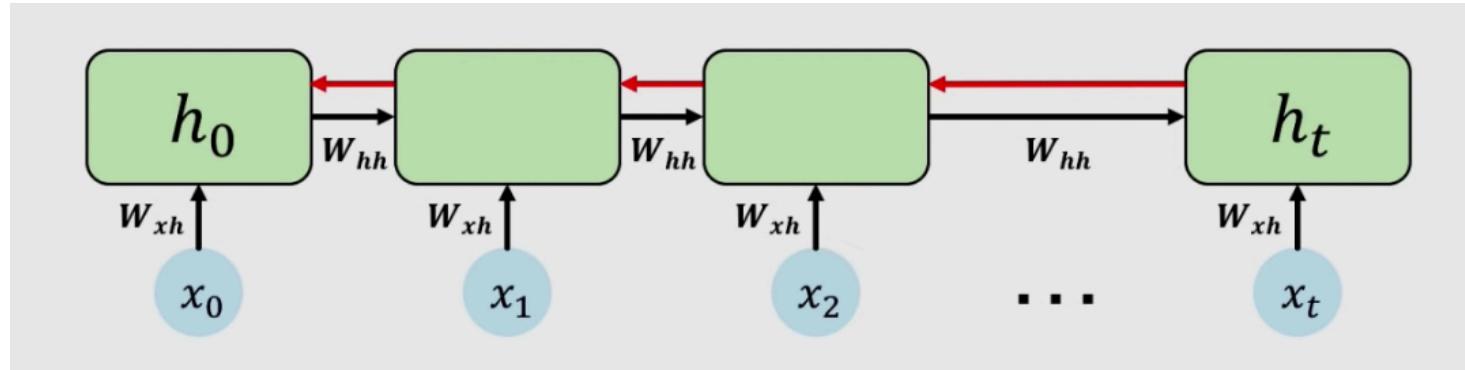


# BACKPROPAGATION THROUGH TIME (BPTT)





BPTT implies a large amount of weight multiplications: if we want to take into account long term memory, networks become quickly very deep and so are subjected to vanishing and exploding gradients problems (see previous lecture)

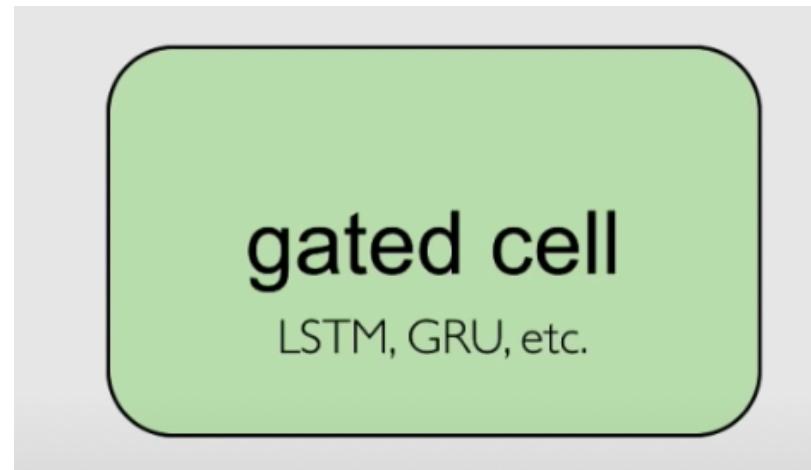


BPTT implies a large amount of weight multiplications: if we want to take into account long term memory, networks become quickly very deep and so are subjected to vanishing and exploding gradients problems (see previous lecture)

ONE CAN USE THE SAME TRICKS THAT WE DISCUSSED  
FOR DEEP CNNs (WEIGHT INITIALISATION, RELU  
ACTIVATION FUNCTION ...)

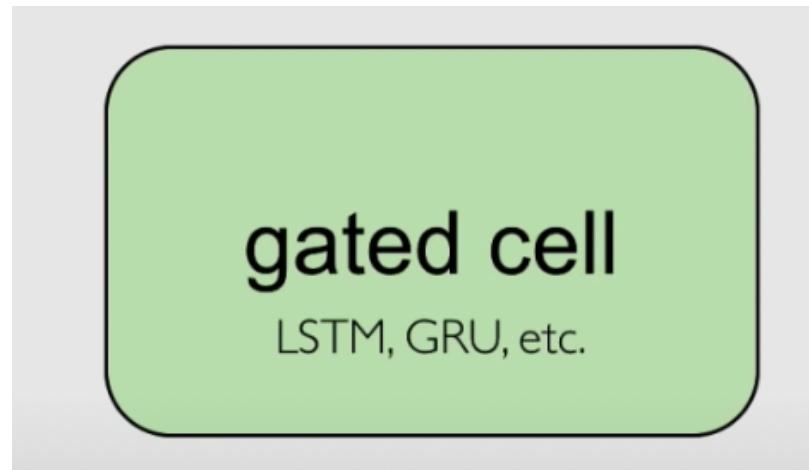
## ANOTHER SPECIFIC SOLUTION: GATED CELLS

Use a more complex recurrent unit with gates to control what information is passed through...



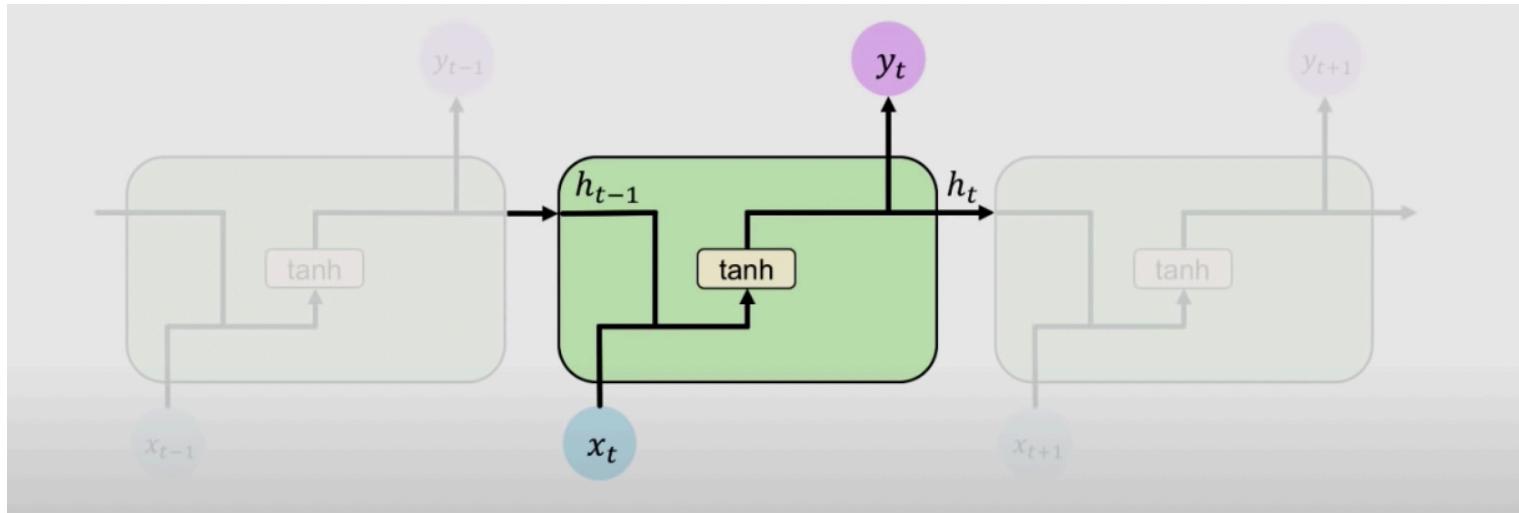
## ANOTHER SPECIFIC SOLUTION: GATED CELLS

Use a more complex recurrent unit with gates to control what information is passed through...

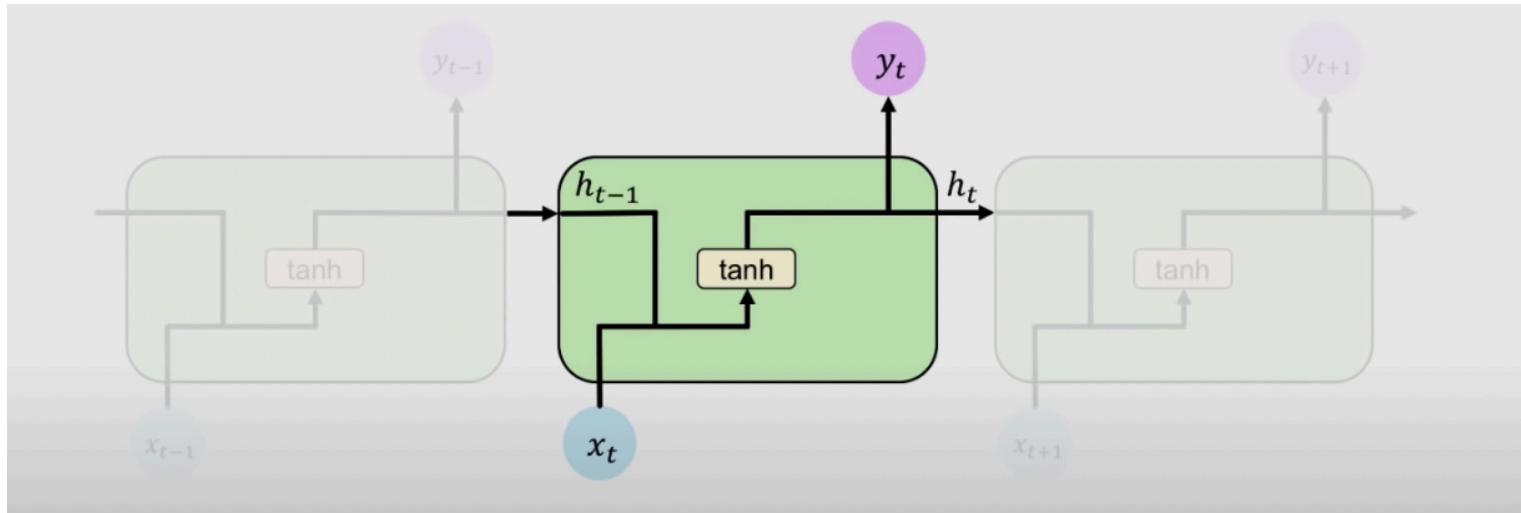


Long Short Term Memory (LSTMs) networks are an example of  
this

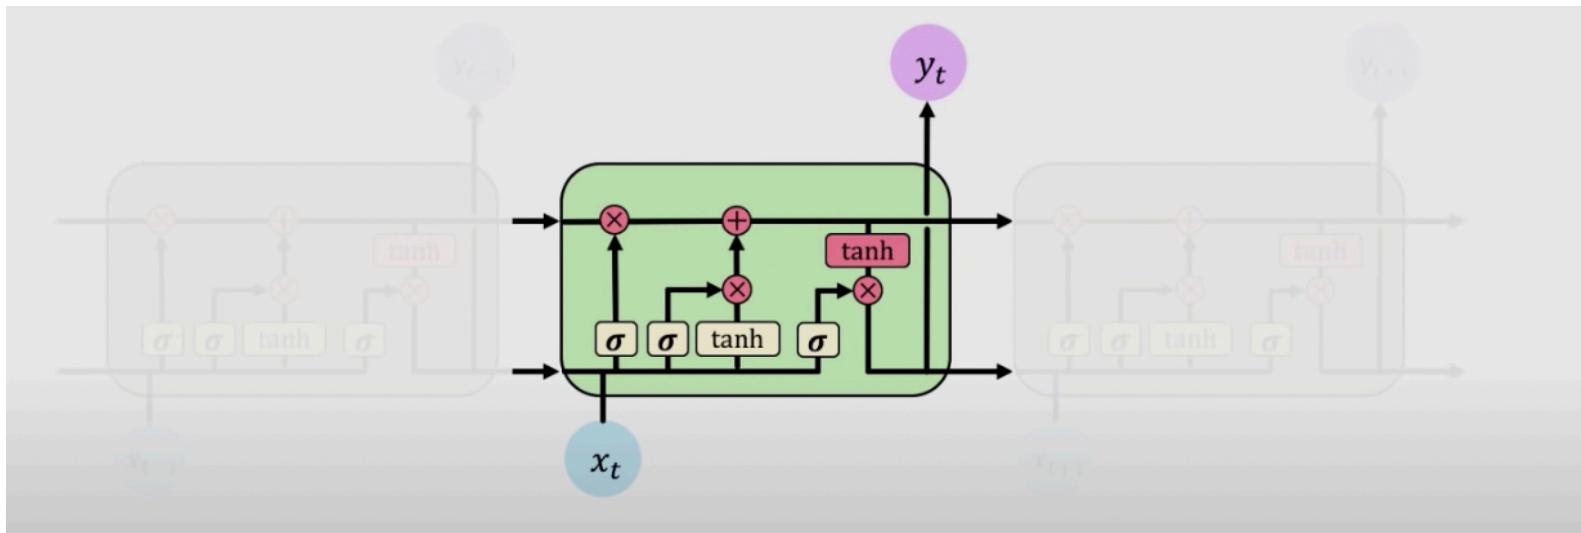
# Standard RNN



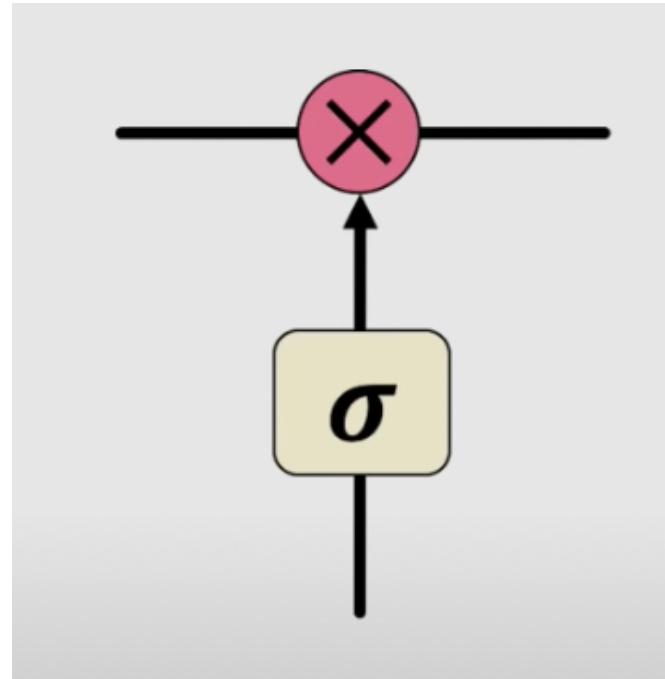
# Standard RNN



# LSTM unit

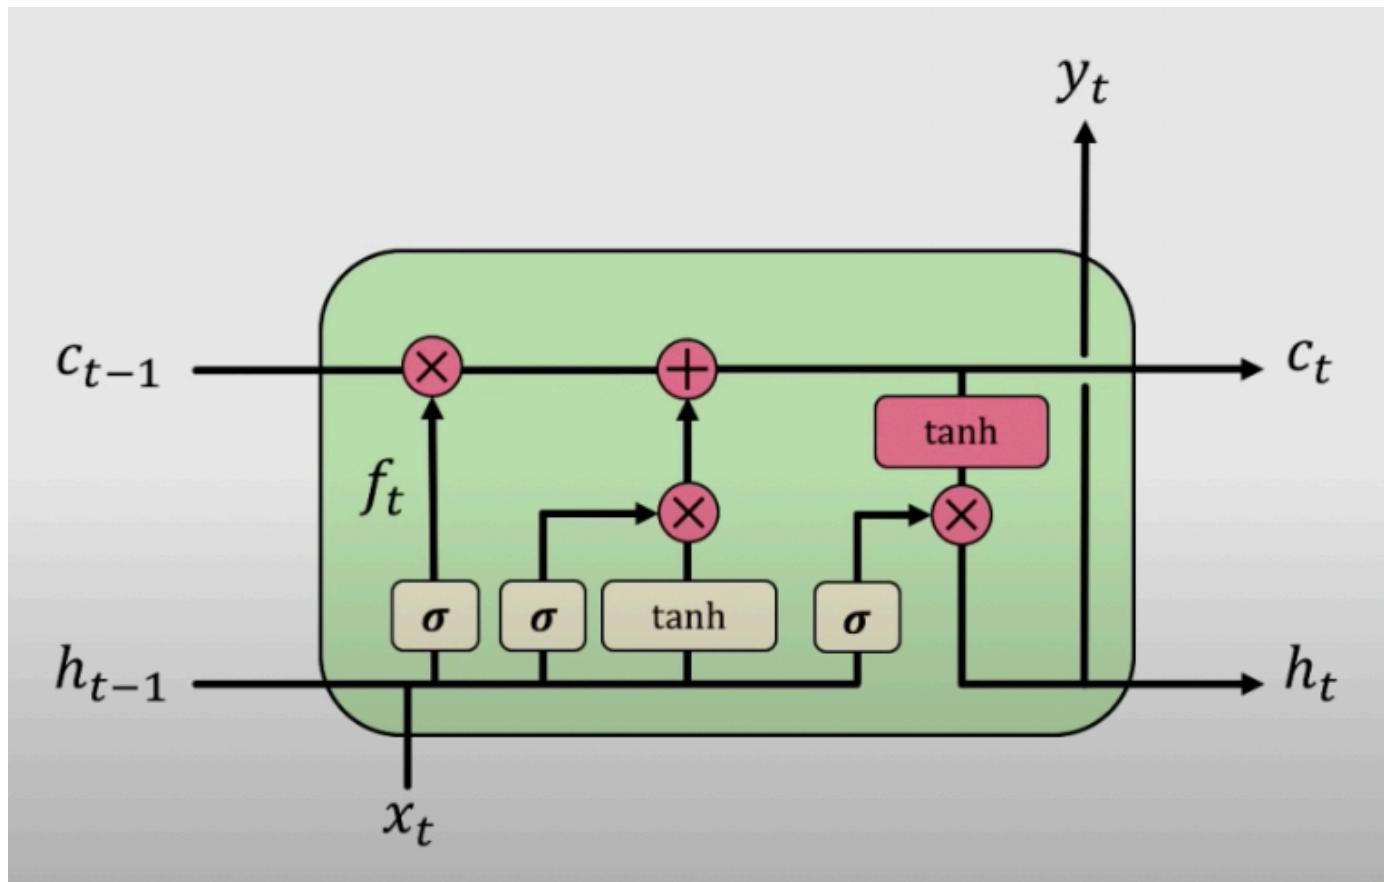


The key building block are the so-called gates



INFORMATION IS PASSED OR NOT WITH A  
COMBINATION OF A SIGMOID NN AND A  
POINTWISE MULTIPLICATION

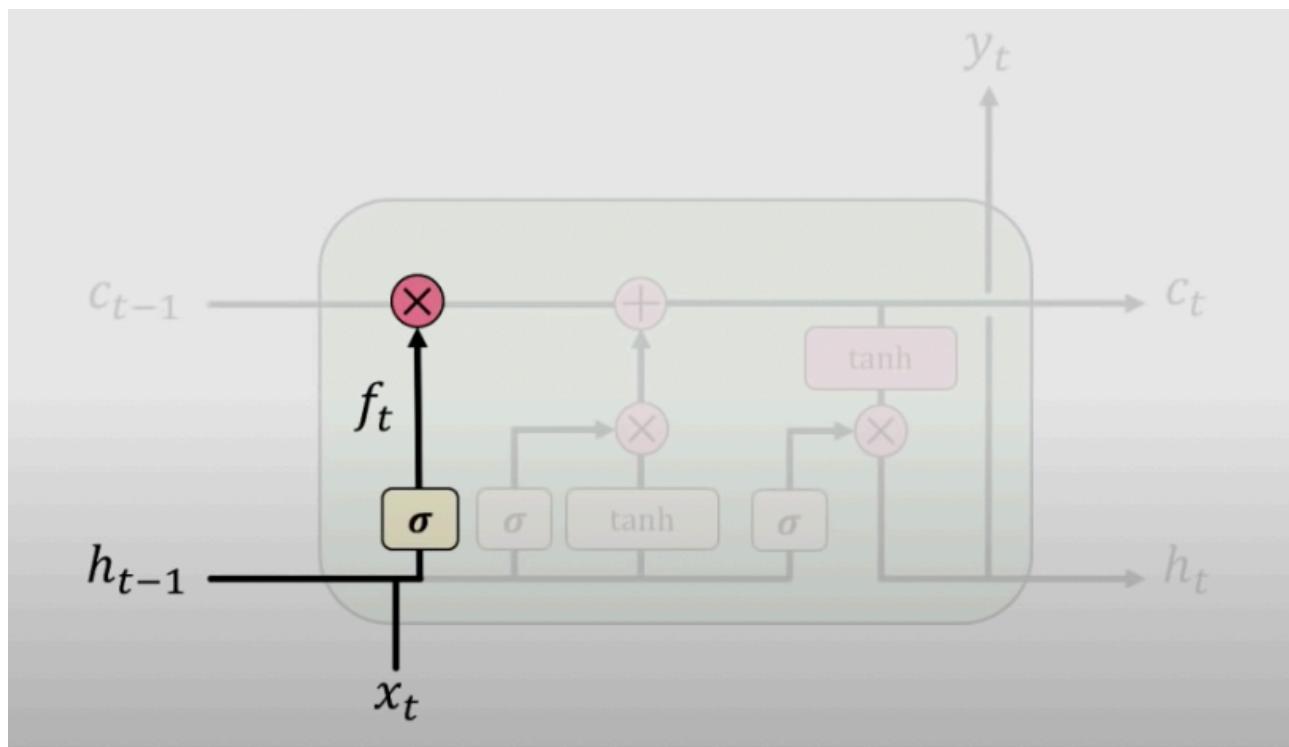
# Long Short Term Memory (LSTMs)



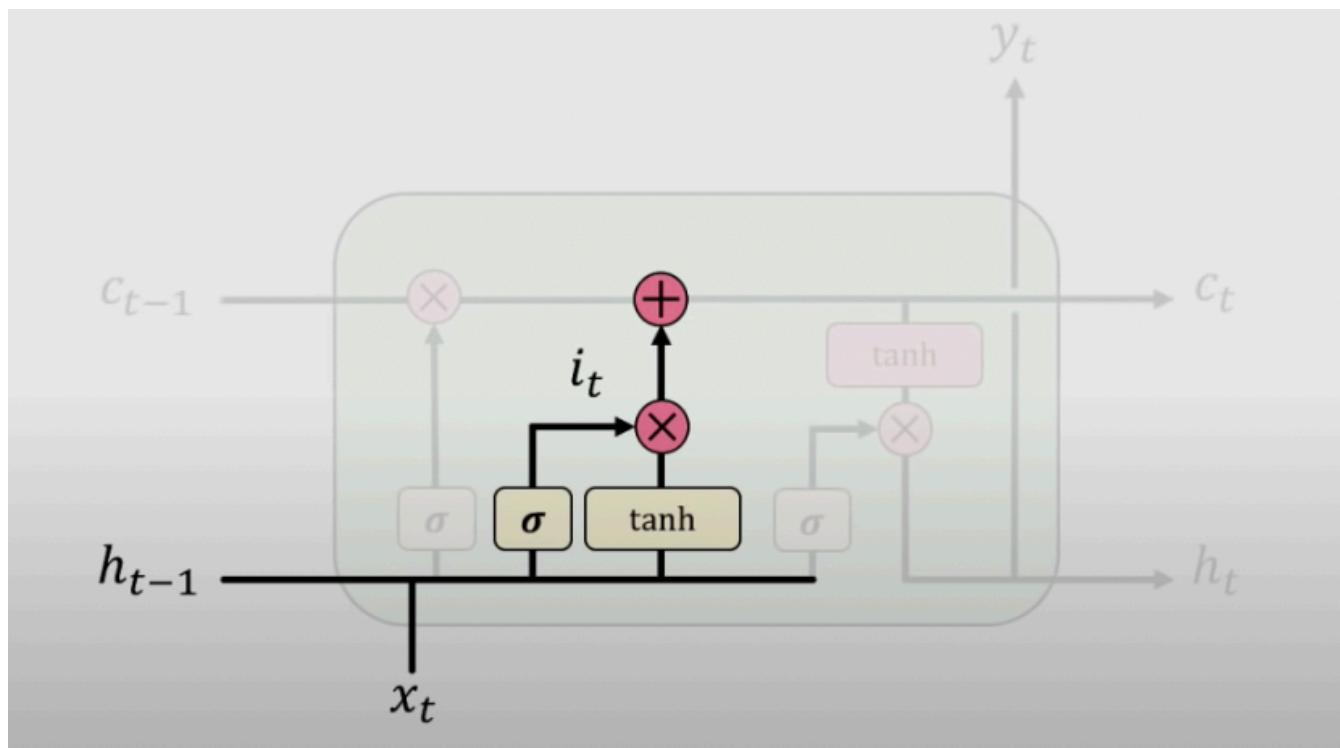
4 main steps:

**1-Forget 2-Store 3-Update 4-Output**

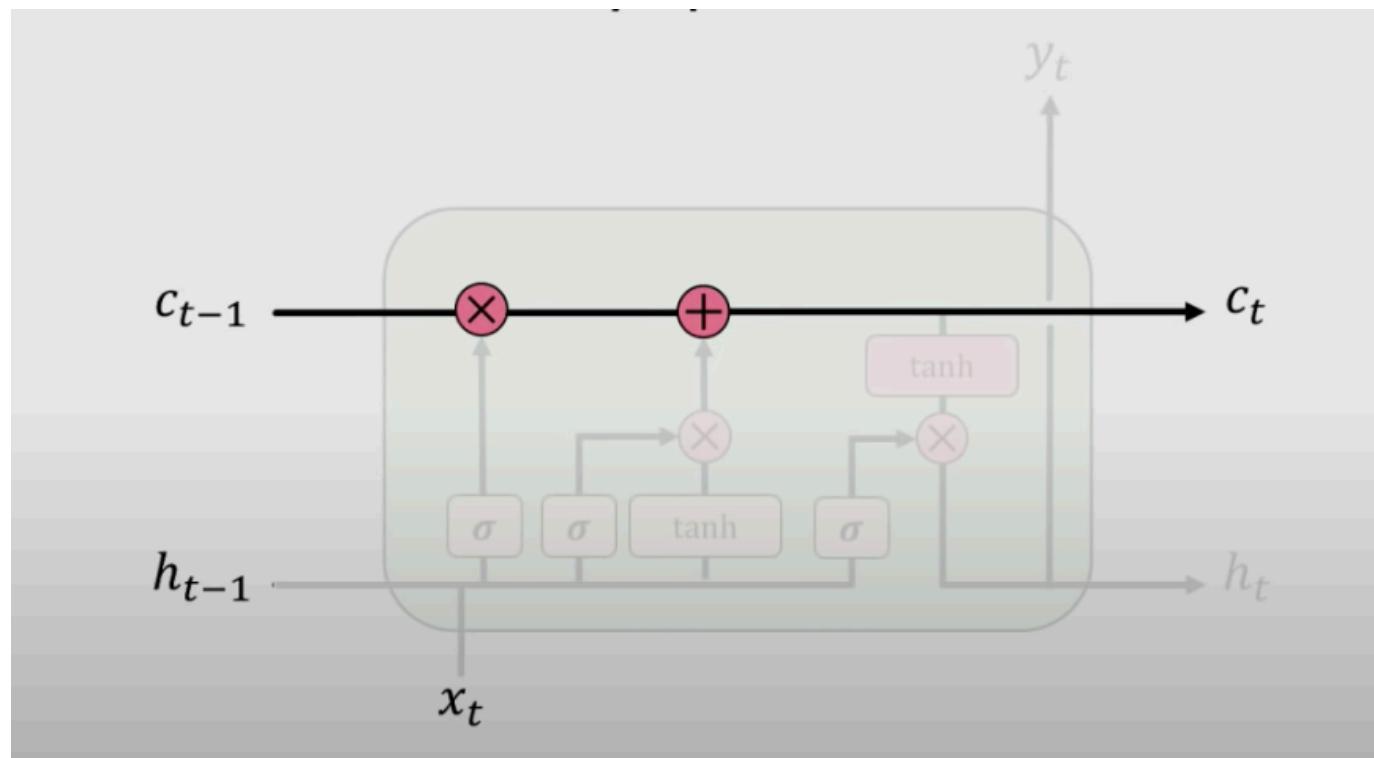
# STEP 1: FORGET



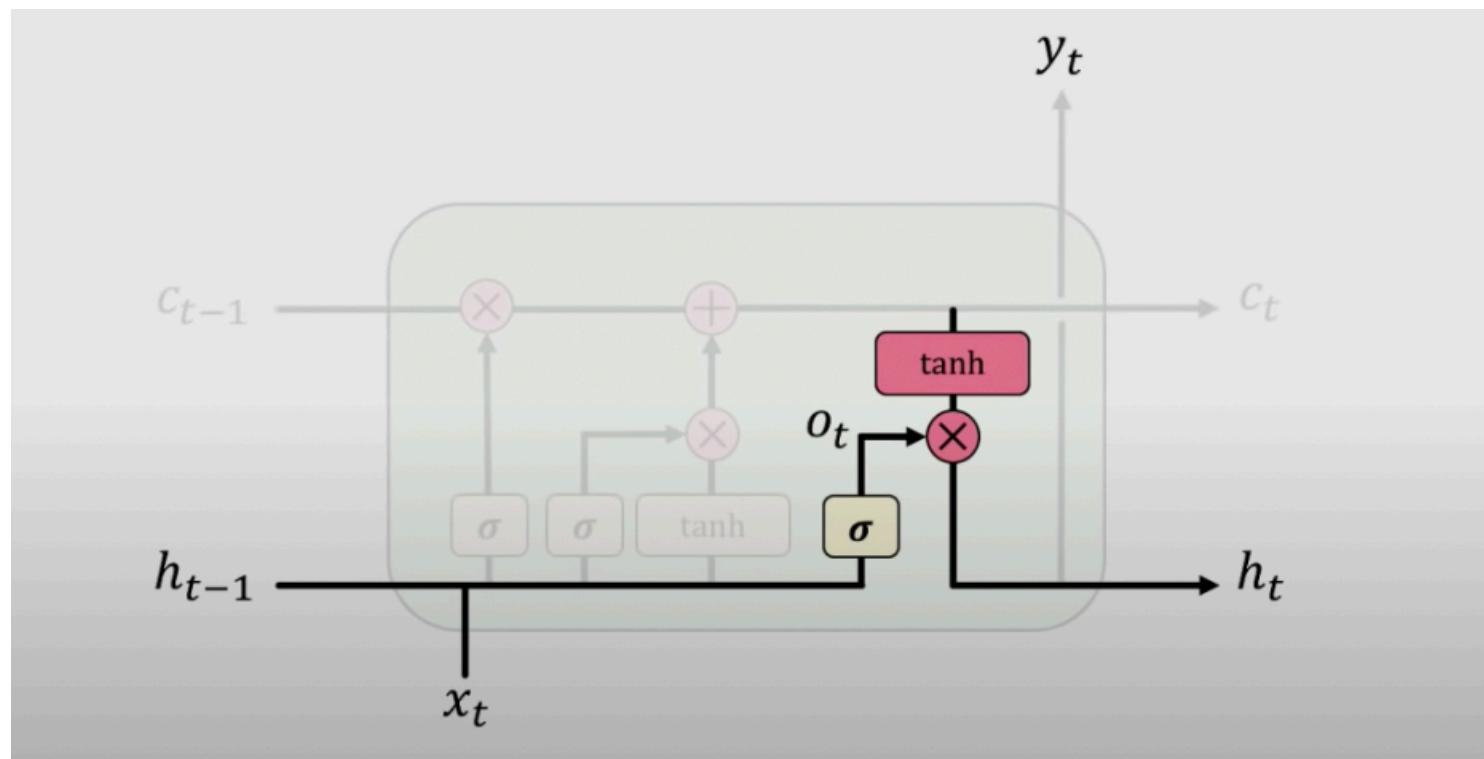
## STEP 2: STORE



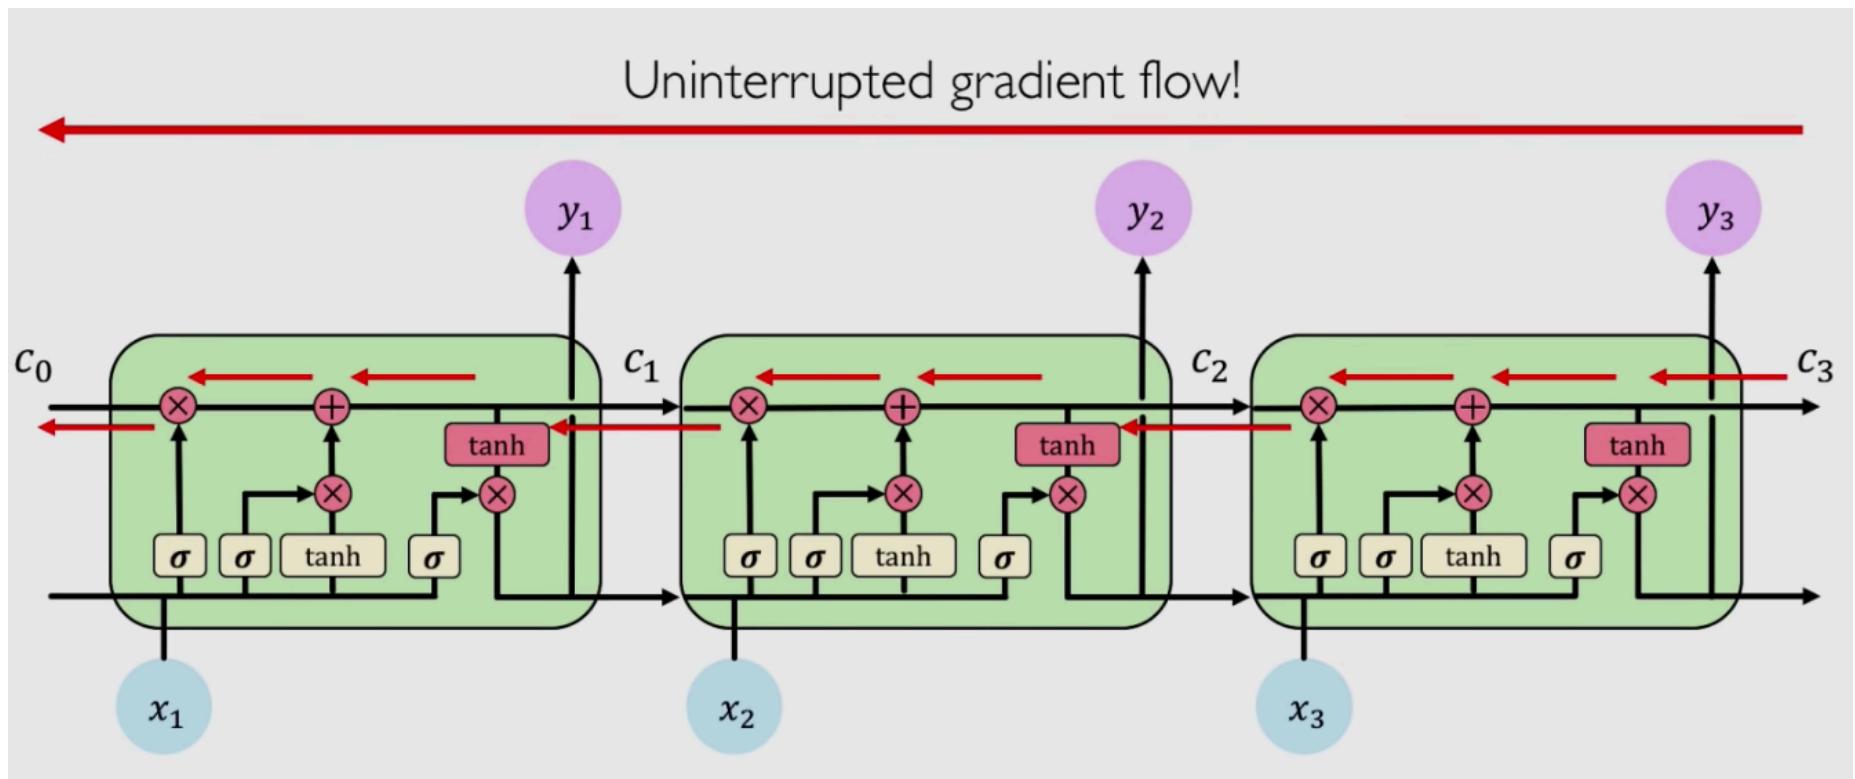
## STEP 3: UPDATE



## STEP 4: OUTPUT



# LSTMs help mitigating the vanishing gradient problem



## **LSTMs KEY CONCEPTS**

**1. Maintain a separate cell state from what is outputted**

**2. Use gates to control the flow of information**

- Forget gates get rid of irrelevant information
- Store relevant information from current input
- Selectively update cell state
- Output gate returns a filtered version of the cell state

**3. Backropagation through time with uninterrupted gradient flow**